
MRNet for Knee Diagnosis

Introduction

Magnetic resonance imaging (MRI) is the most preferred used method by clinicians to diagnose knee injuries but it takes long time to diagnose them and is subject to error so Deep learning methods are developed to help in the interpretation of the medical images . one of these methods is the MRNet model which cares about detecting general abnormalities and special abnormalities like anterior cruciate ligament (ACL) tears and meniscal tears.



Dataset

We have **1250 MRI exams** in which **1130 exams** for training and **120** for validation so we classified them as following:

- **Training**
90% of training data (1017 exams)
- **Validation**
10% of training data (113 exams)
- **Test**
whole Validation data (120 exams)

Exams setup

Each MRI exam consists of S slices each one is of size $(256 \times 256 \times 3)$. Exams of patients are passed through three different series (Axial, Sagittal and Coronal) to check the probability of detecting the anomaly (abnormal, ACL and Meniscal) from each series.

Model

It mainly consists of three parts :

→ **Feature extractor**

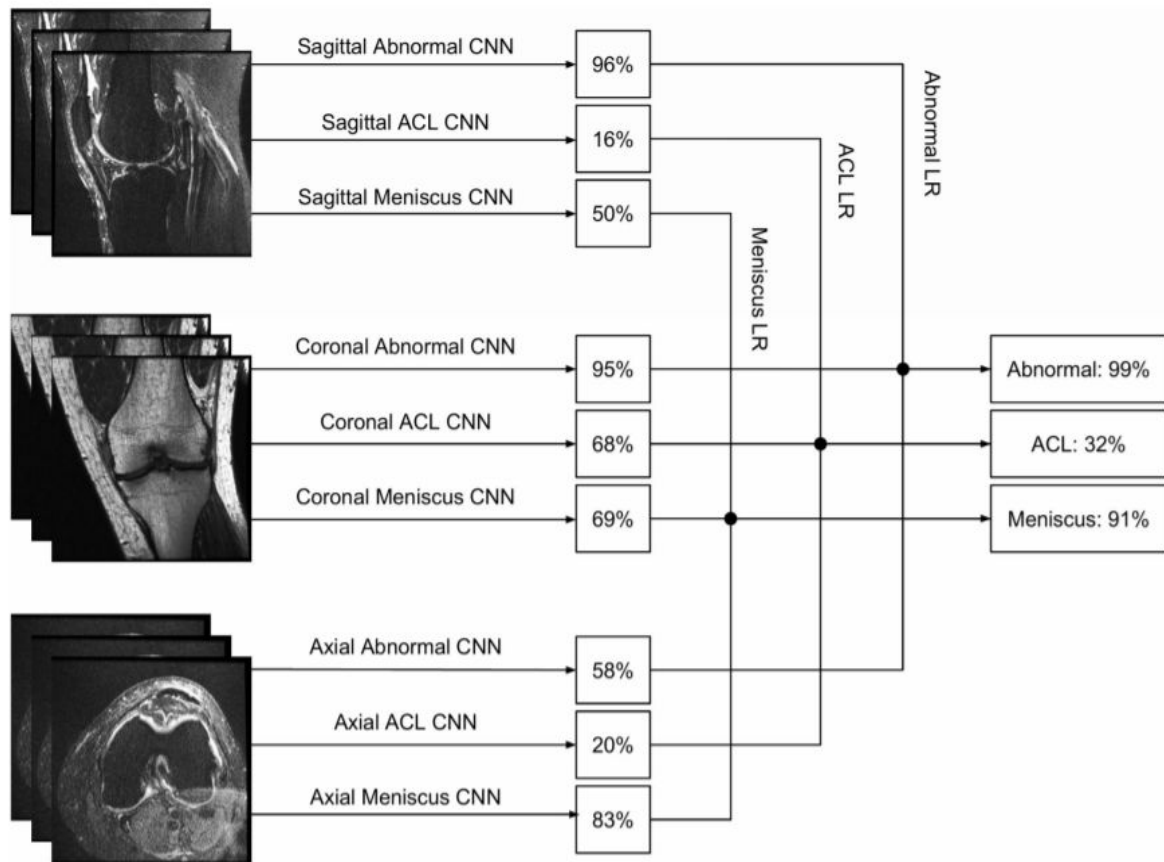
For training, it takes the first slice from a certain series to specify the features of each anomaly according to a specific series. So, we have 9 total.

→ **Classifier**

To detect the probability of each anomaly according to its series. Also, 9 total.

→ **Regressor**

It finds the probability of detecting the anomaly by taking the predictions of three classifiers for the same anomaly from the three different series. 3 total.



Feature extractors models

- We used three different CNN models for training the feature extractors :

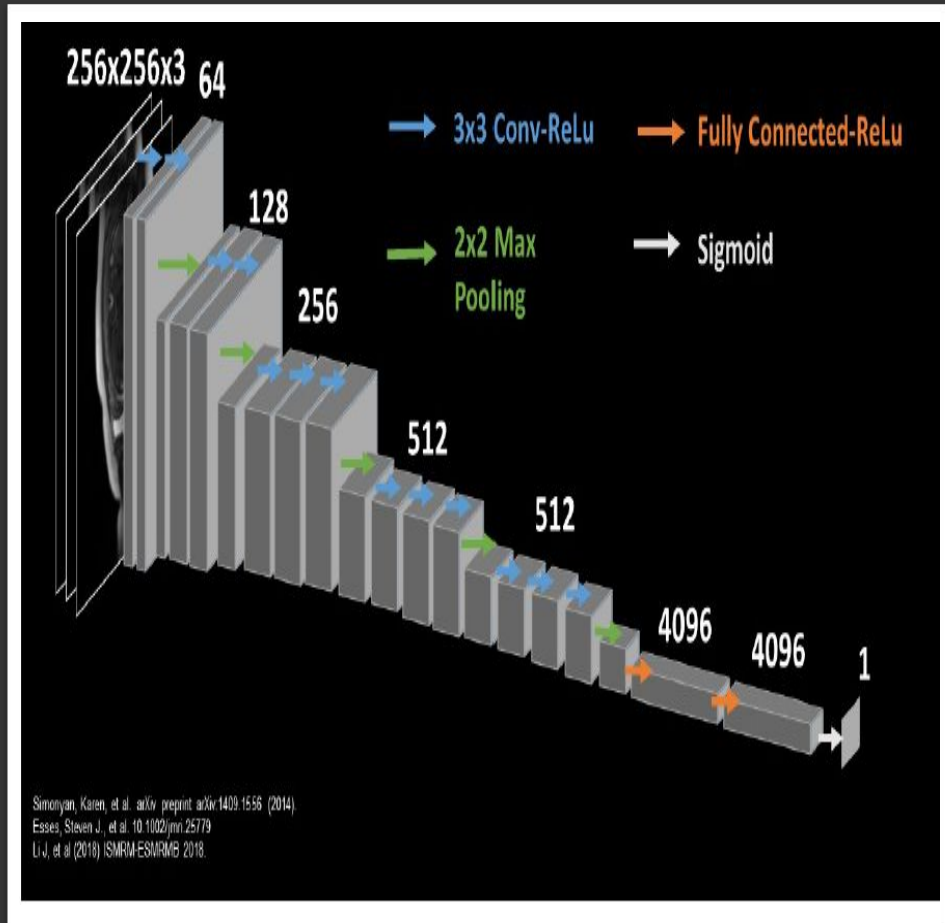
VGG

RESNet

Inception V3

VGG

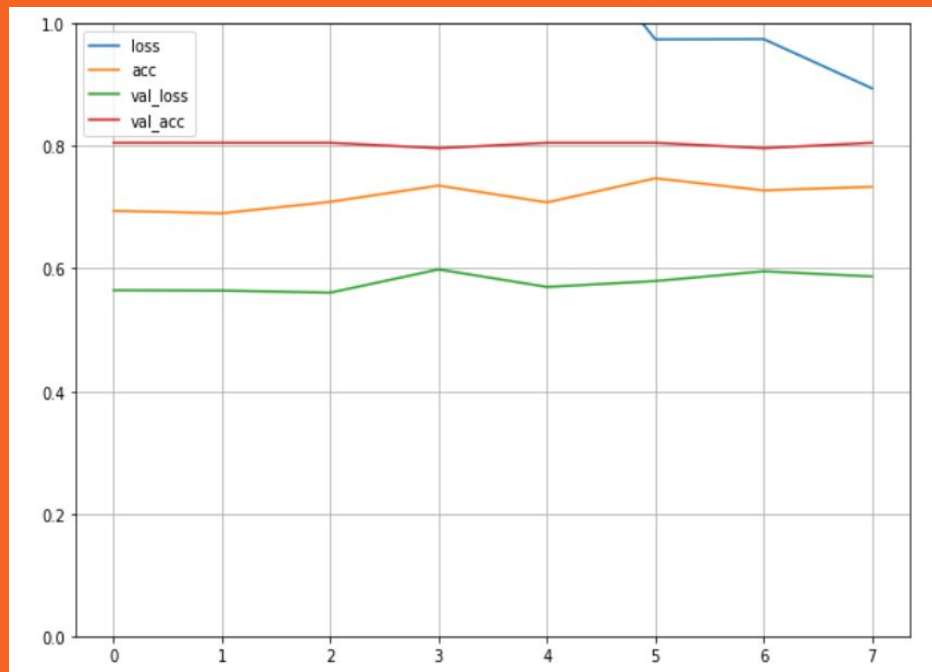
- It is a simple neural network architecture of 16 layers with simple hyperparameters. All conv layers are 3X3 filters and same padding . for all max pooling layers we used 2 X 2 filters with stride of 2



Training Extractors

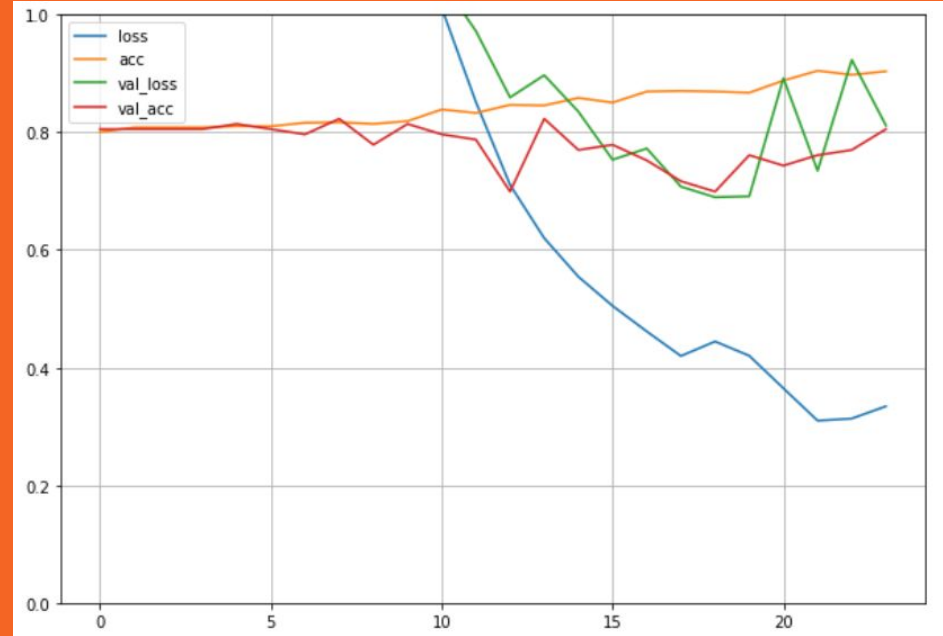
Extractor 1

By using low learning rate (10^{-6}) but the loss decayed slowly.



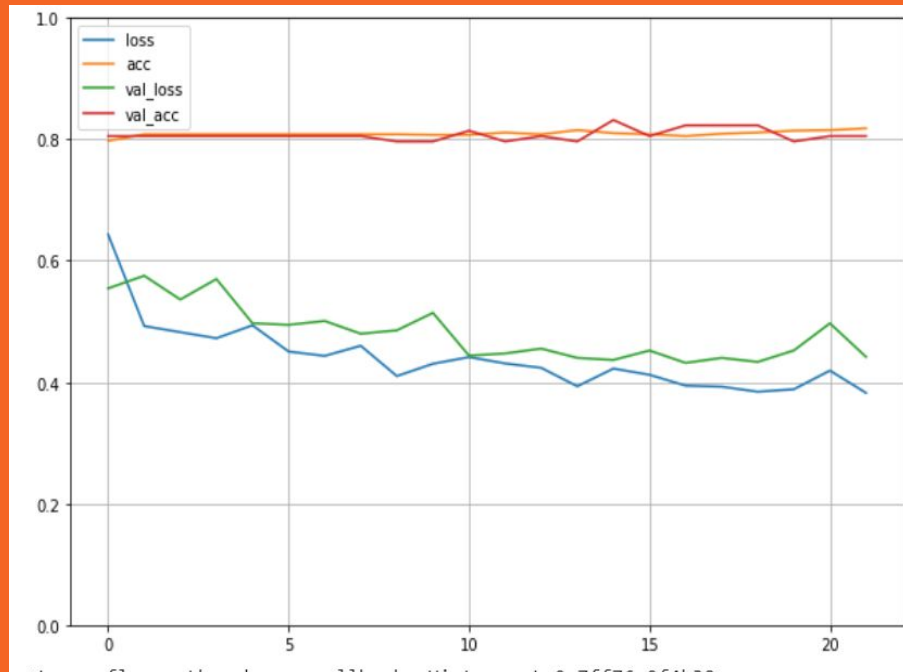
Extractor 2

Using regularization term = 0.01 but it increases the validation loss



Extractor 3

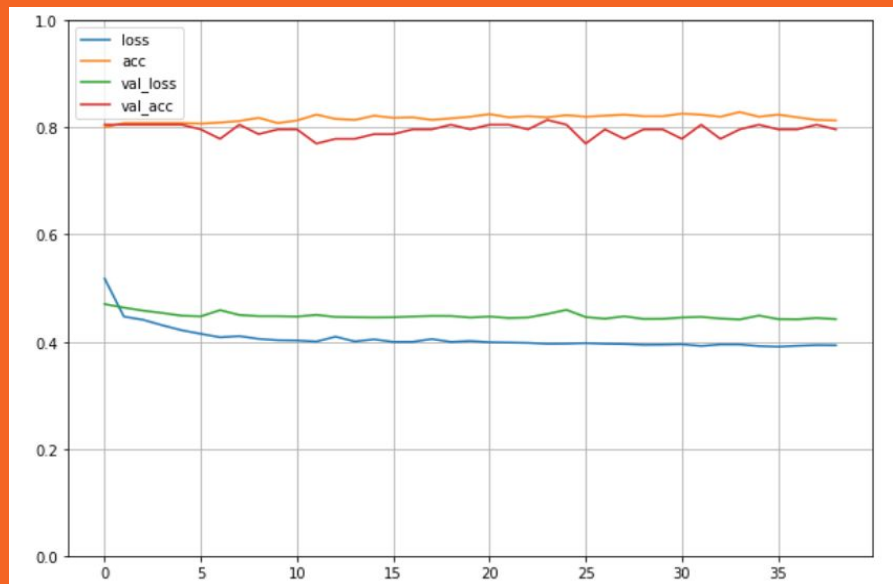
It is the final model for extractor with learning rate (10^{-4}) and use some dropout layers



Training Classifiers

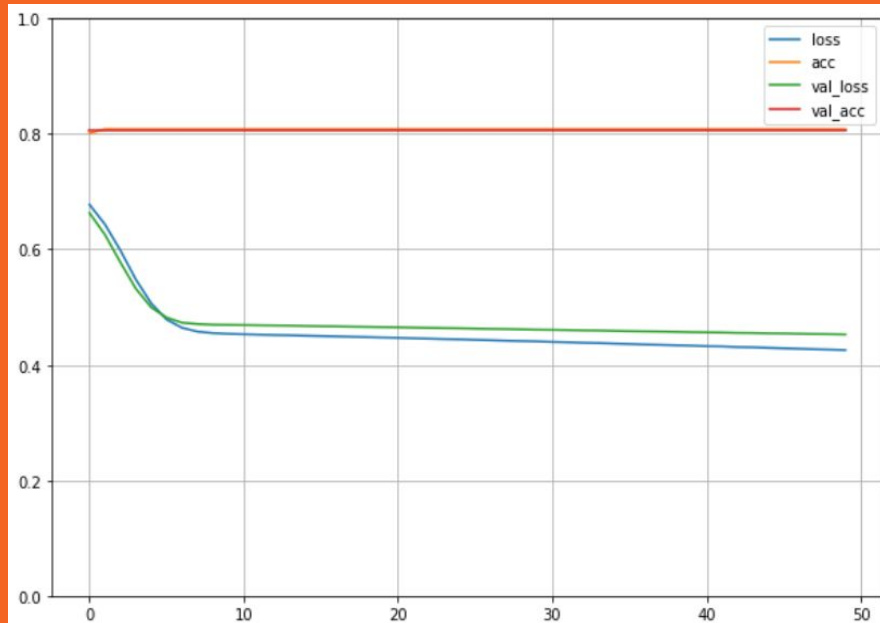
Classifier 1

Using many neurons in dense layers (4096)



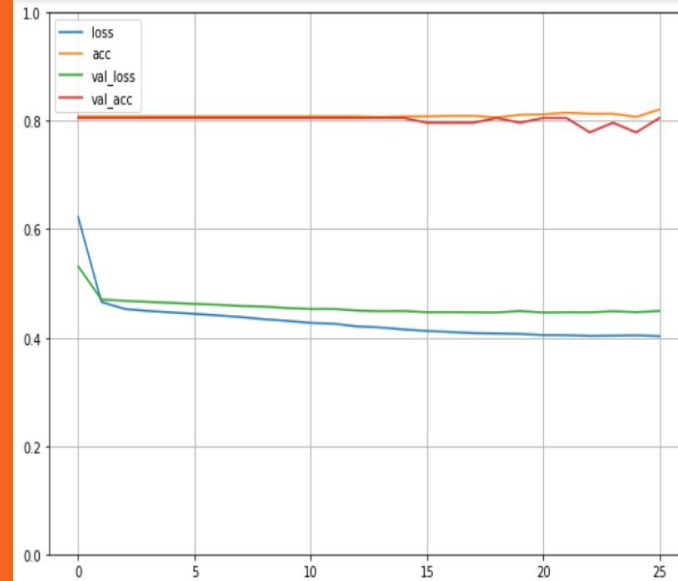
Classifier 2

Using less neurons in dense layers with high learning rate (0.01)



Classifier 3

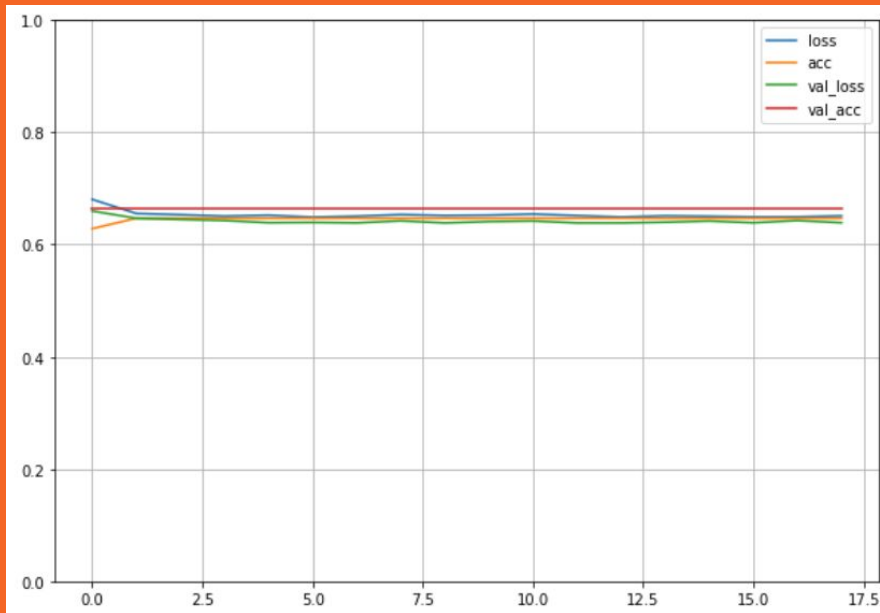
It is the final model for classifier with learning rate (10^{-4}) and use (1024 neurons in first dense layer and 512 in the second one).



Training Regressors

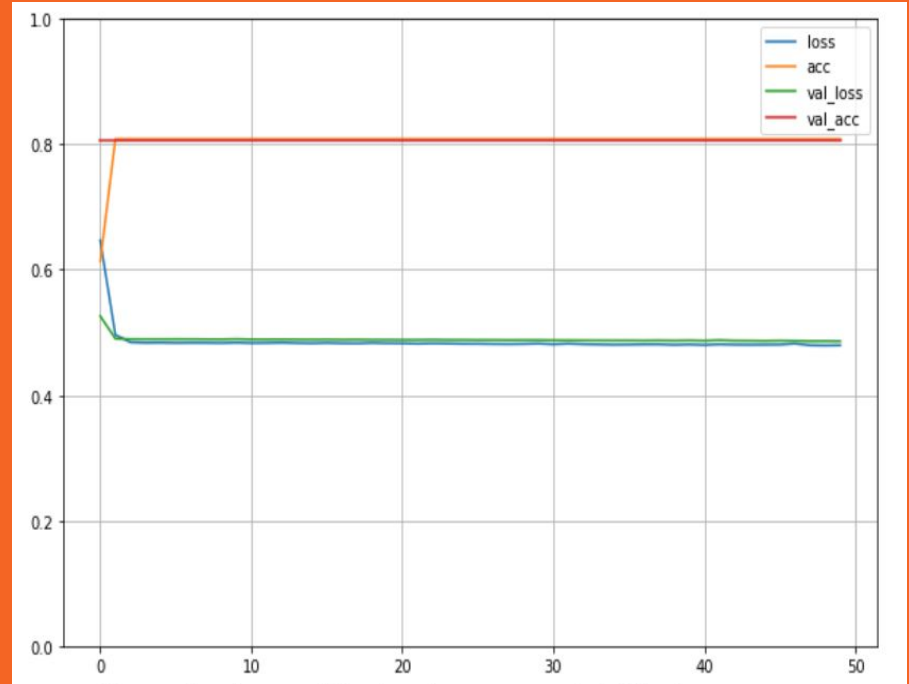
Regressor 1

Using low learning rate (10^{-4})



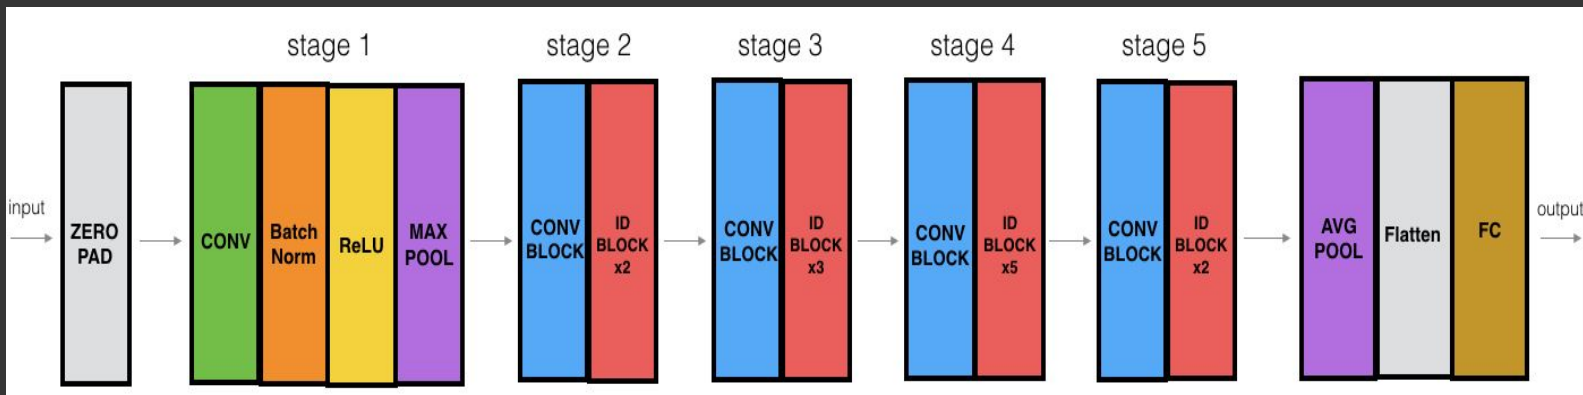
Regressor 2

It is the executed model for regressor by
Using higher learning rate (10^{-2})



RESNet

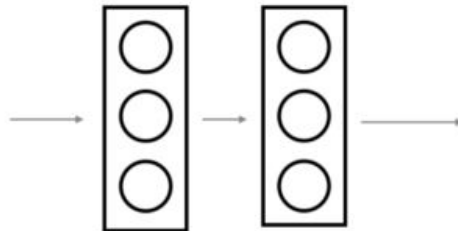
- It introduces the idea of training a very deep neural networks using the residual blocks that allows us to take the activation from a layer and feed it to a further layer in the network.



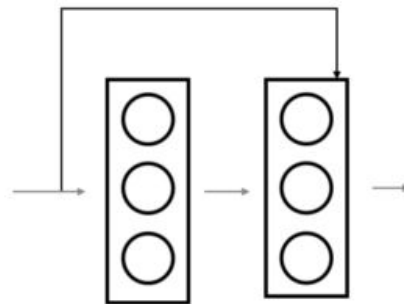
Idea

- When the problem is more complex we need more complex network to solve
- deep networks suffer from overfitting because it may get harder to learn anything in some layers
- In this situation the residual blocks allows the layer to be skipped by feeding the output of the previous layer instead of overfitting

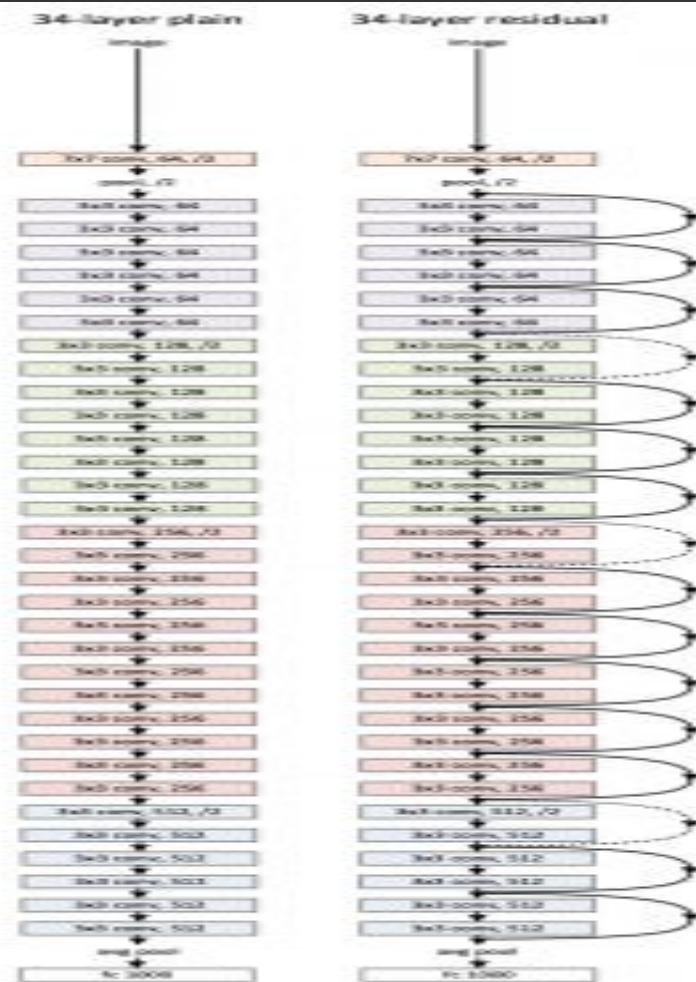
without skip connection



with skip connection

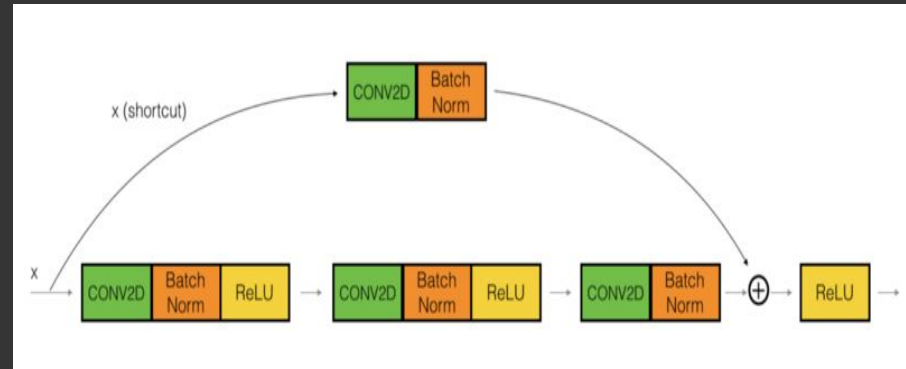
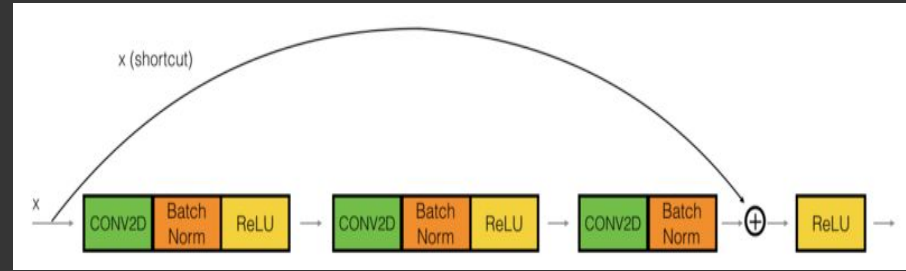


This transforms the regular network to this new shape with the skip connections



Residual blocks

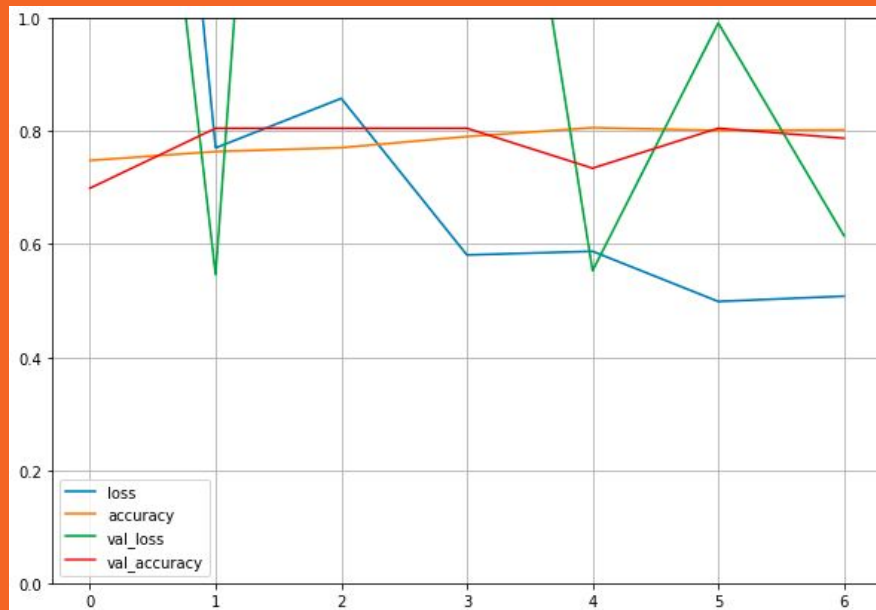
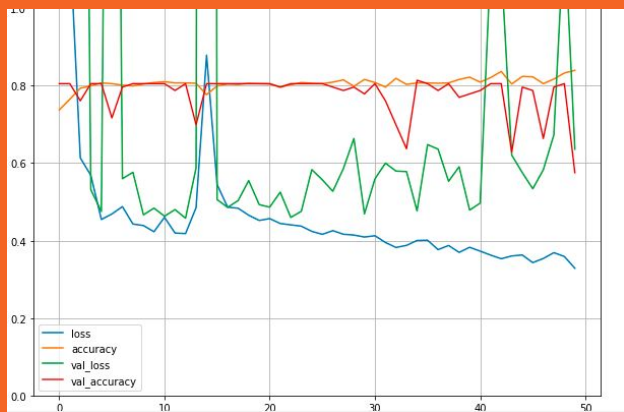
- The skipping operation is implemented via adding the i/p of the block to the o/p of the same block before performing the activation
- If the dimensions are compatible we will add them, if not we will need to pass the i/p via conv layer to make them add compatible
- That's why we have 2 types of block identical, conv blocks



Training Extractors

Extractor 1

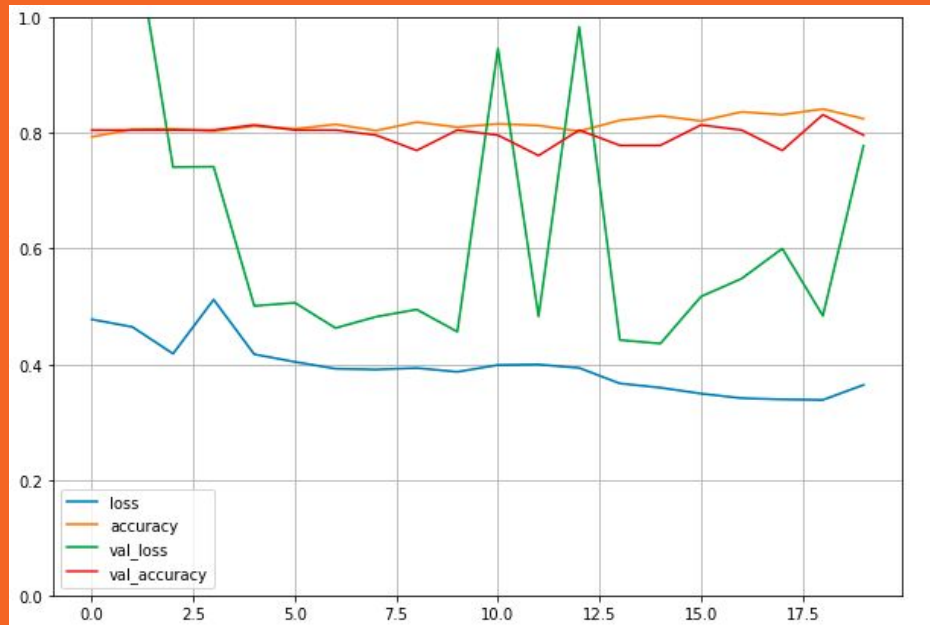
First attempts were too bad without any modifications



Extractor 2

Adjusting the params and the dense layer

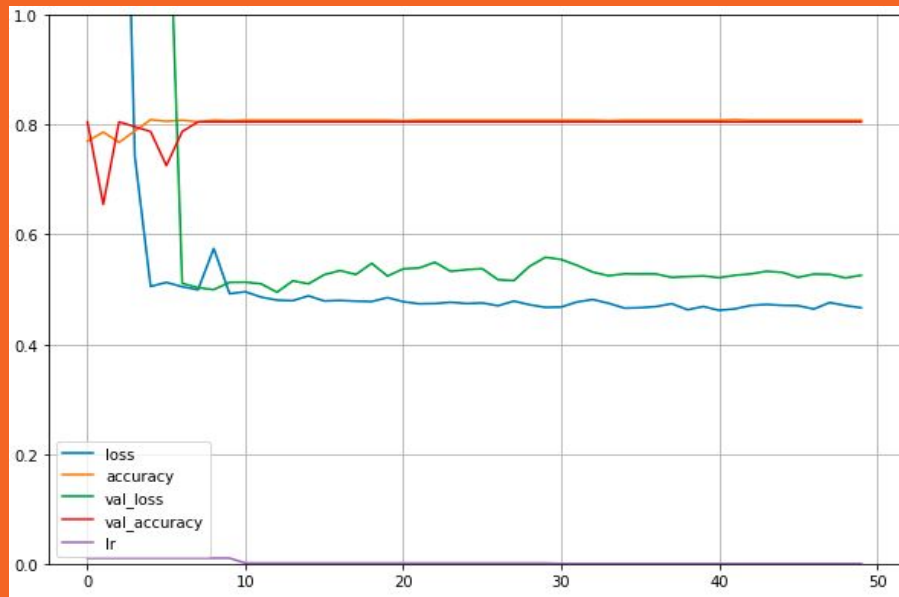
it was better but seems to over fit



Extractor 3

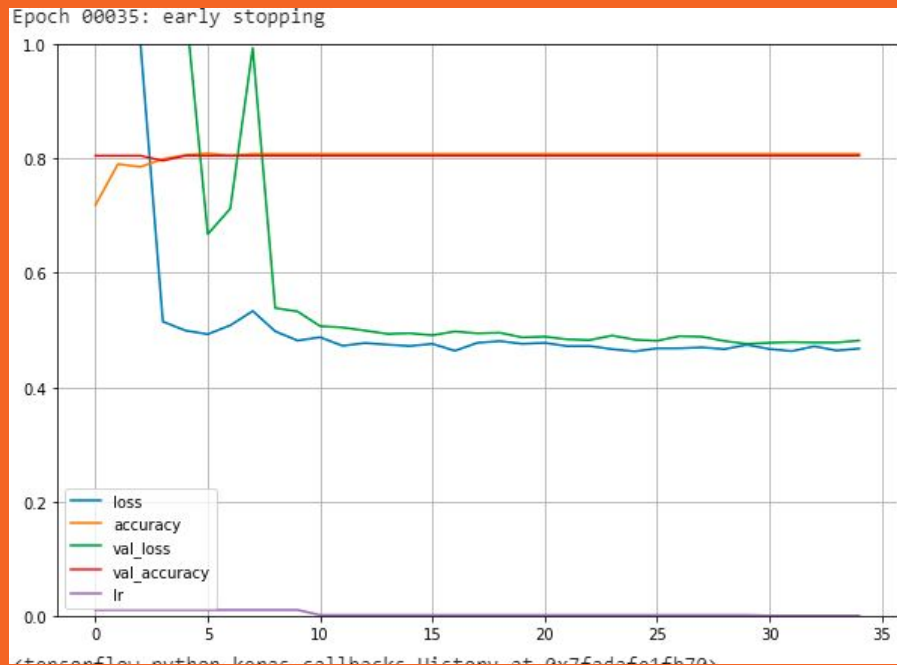
Adding learning rate scheduler call-back to start with 0.01 then 0.001 then if we continue it will be 0.0001

And adding a dropout layer



Extractor 4

Now adding the early stop call-back



Training Classifiers

Classifier 1

With the classifier the same params from the extractor we just fine

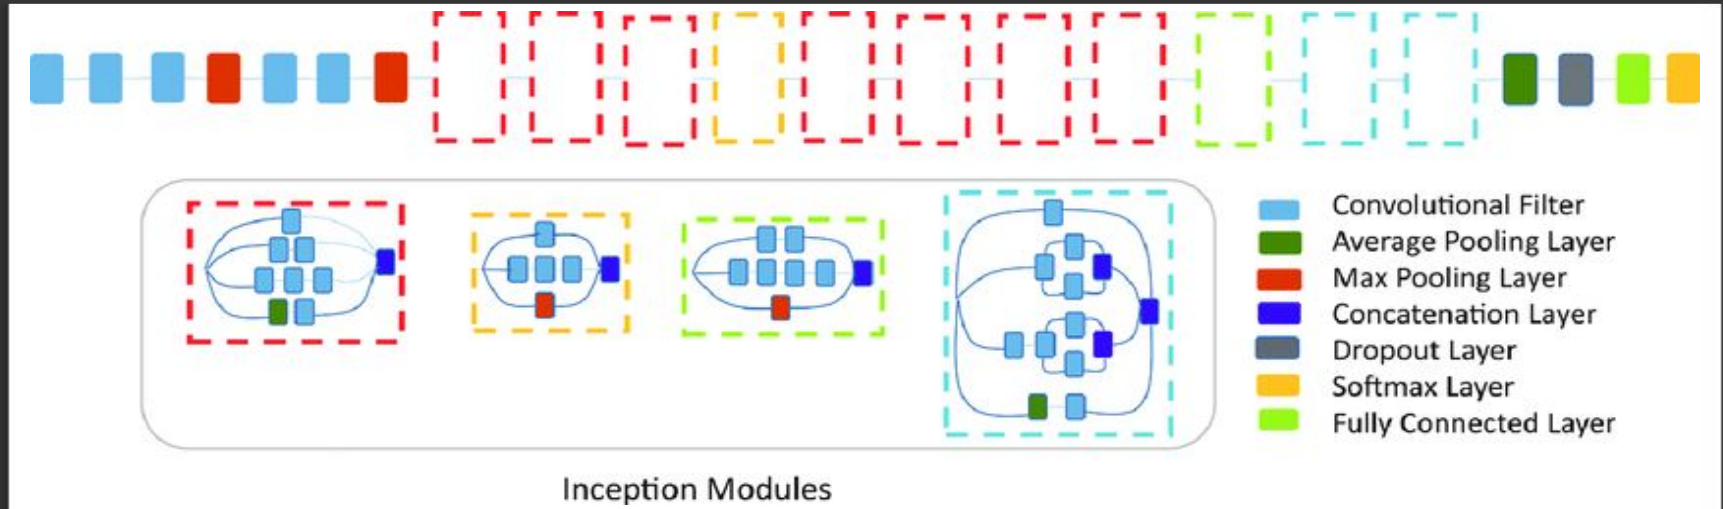
Training Regressors

Regressor 1

Using the same previous params

Inception V3

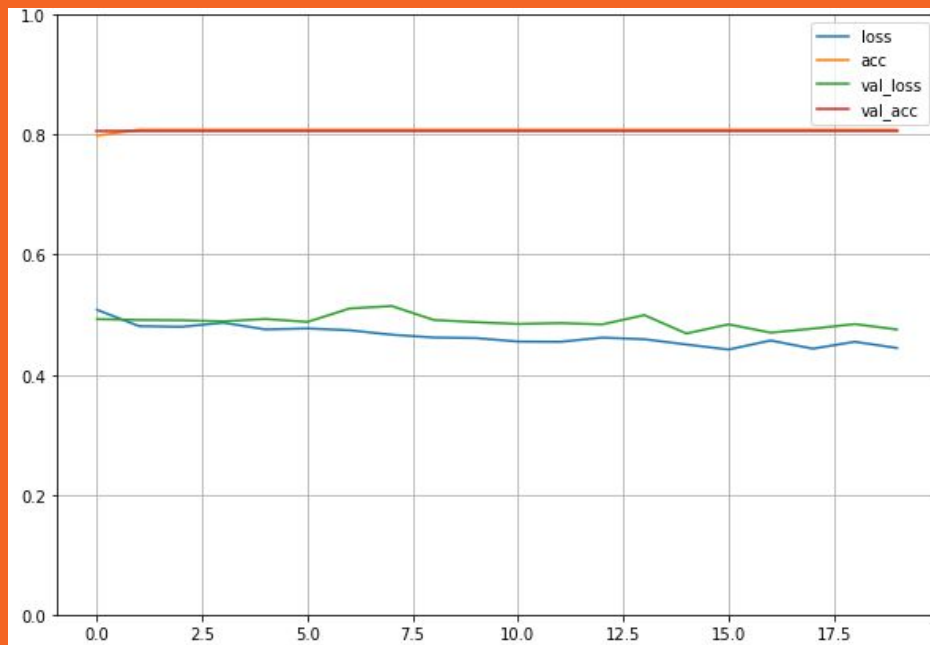
- Instead of using all filters of the same size for a specific layer, We use different sizes (eg: 3×3 , 1×1 , ...) of filters for a layer. It also helps in reducing the computational cost by using 1×1 filters.



Training Extractors

Extractor

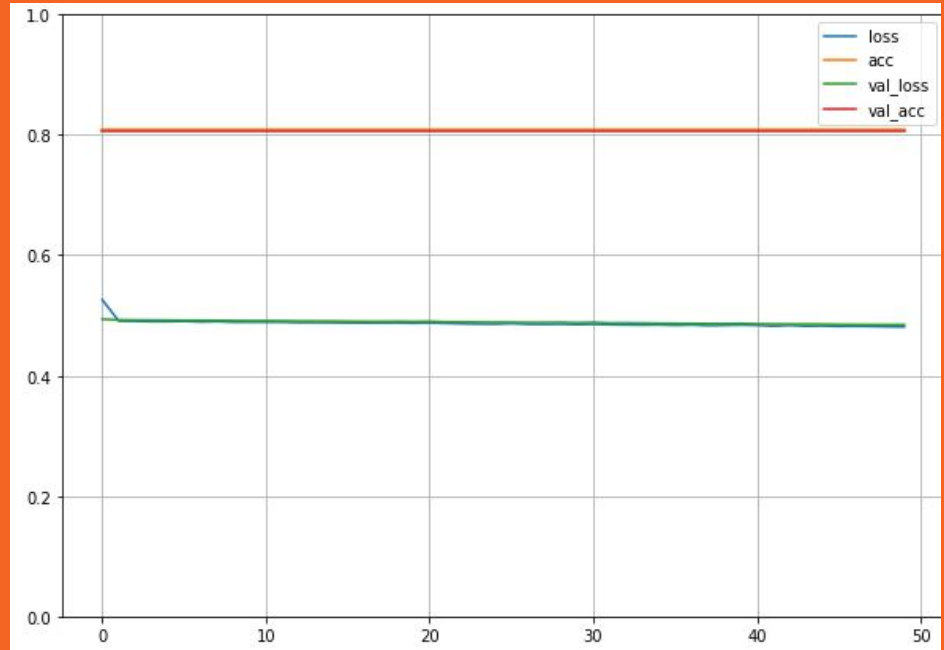
Decreasing the learning rate from 10^{-3} to 10^{-5} we now can see the changing in loss as it's decreasing but before it was approaching straight line.



Training Classifiers

Classifier

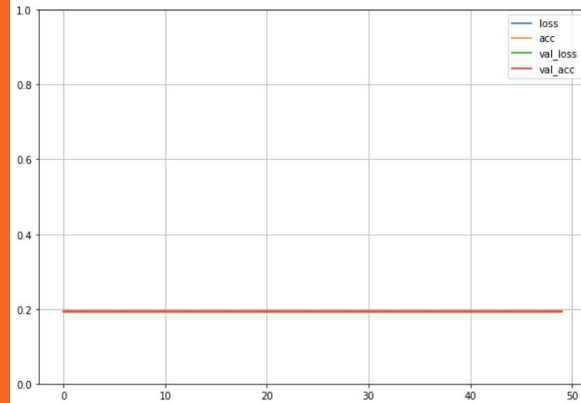
Using learning rate 10^{-5} there is a little decreasing in loss but this can be changed by using different learning rate



Training Regressors

Regressor 1

Low Learning rate 10^{-5} for abnormal regresor
- failed trial

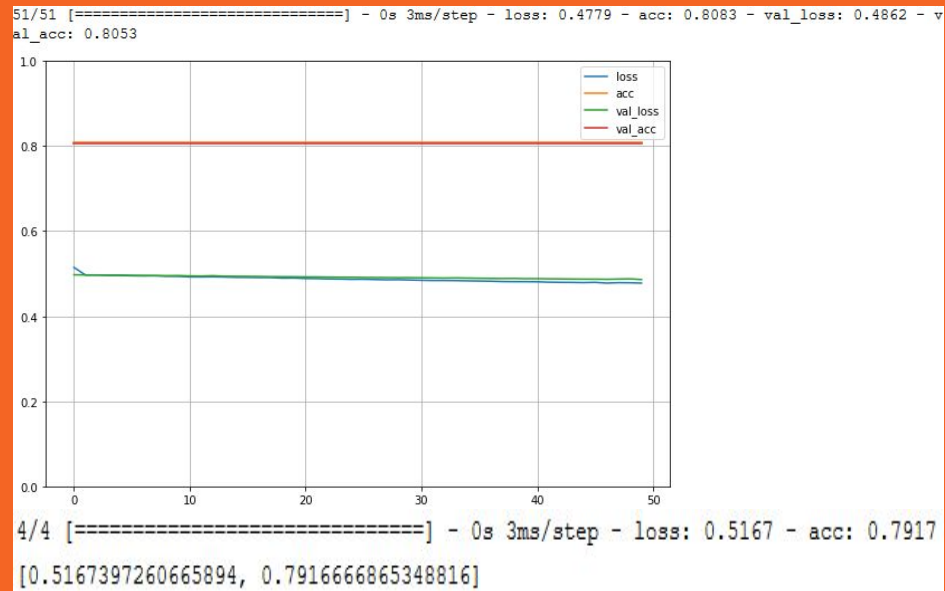


```
processing.test_regressor(axial_extractor,sagittal_extractor,coronal_extractor, processing.inception, processing.abnormal)
```

```
4/4 [=====] - 0s 2ms/step - loss: 1.0370 - acc: 0.2083  
[1.0369518995285034, 0.2083333283662796]
```

Regressor 2

Abnormal regressor after using learning rate 10^{-2}



Statistics



Regressors

We judge the model by two methods

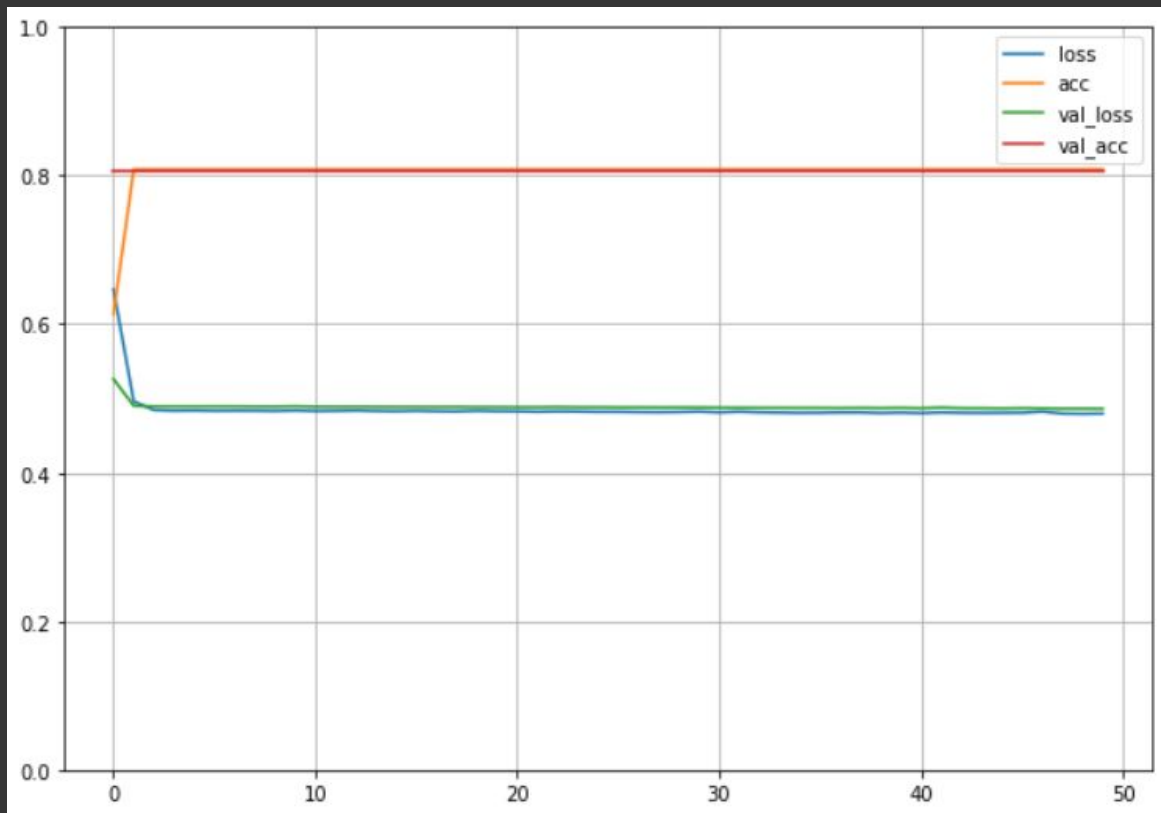
- Loss
- Accuracy

VGG

- **Abnormal**

Loss : 0.5356

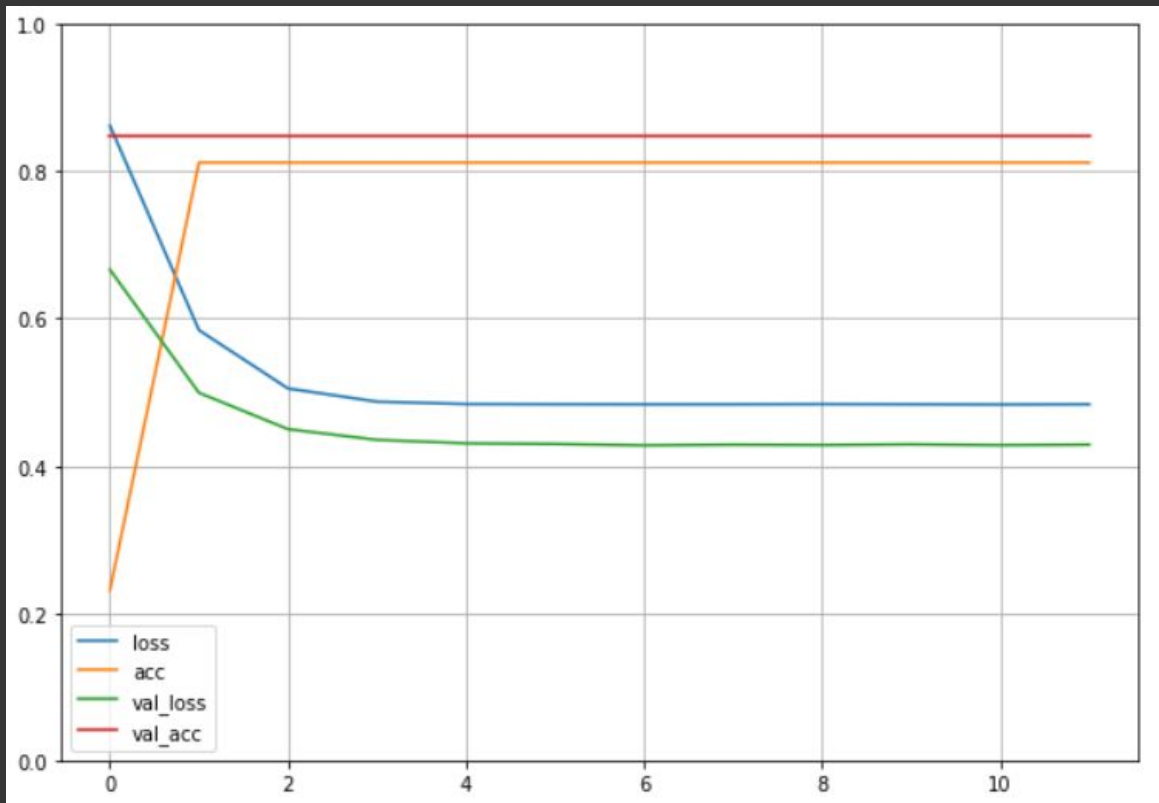
Accuracy : 0.7917



VGG

- **ACL**

Loss : 0.6903
Accuracy : 0.5500

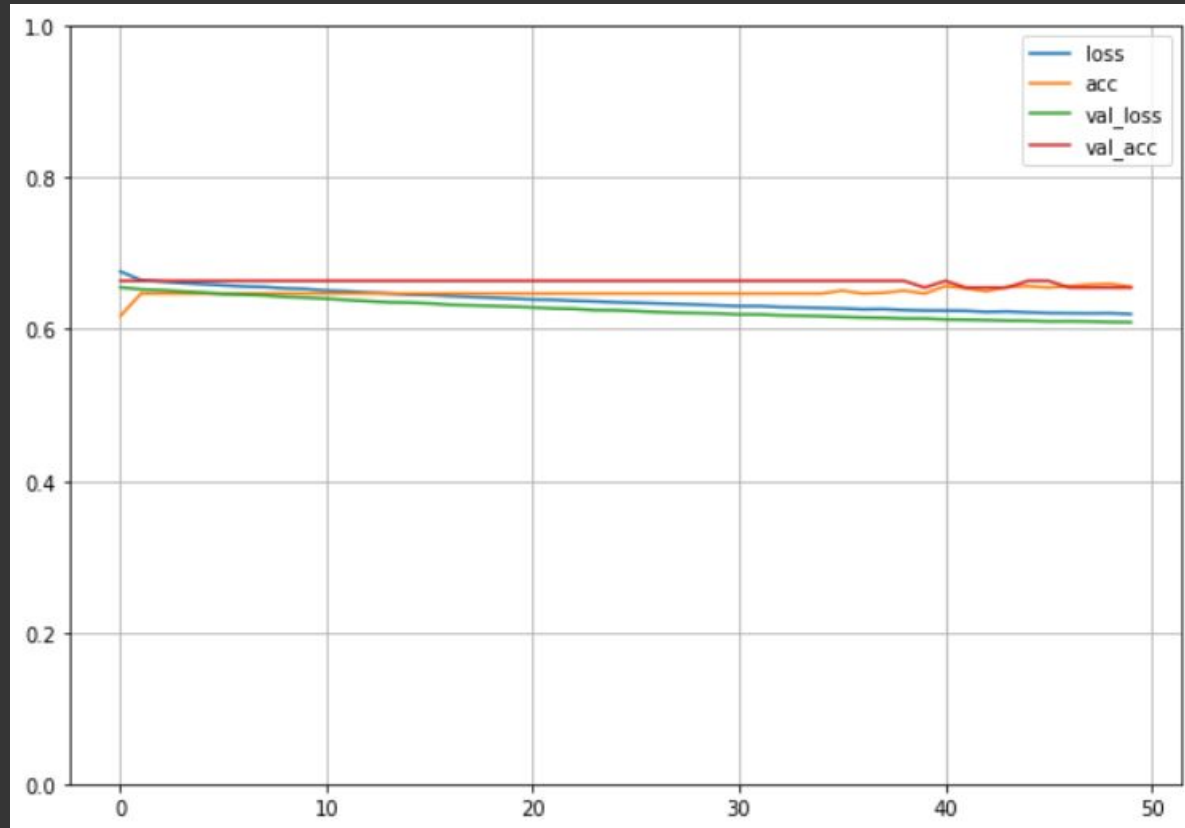


VGG

- **Meniscal**

Loss : 0.6982

Accuracy : 0.5667

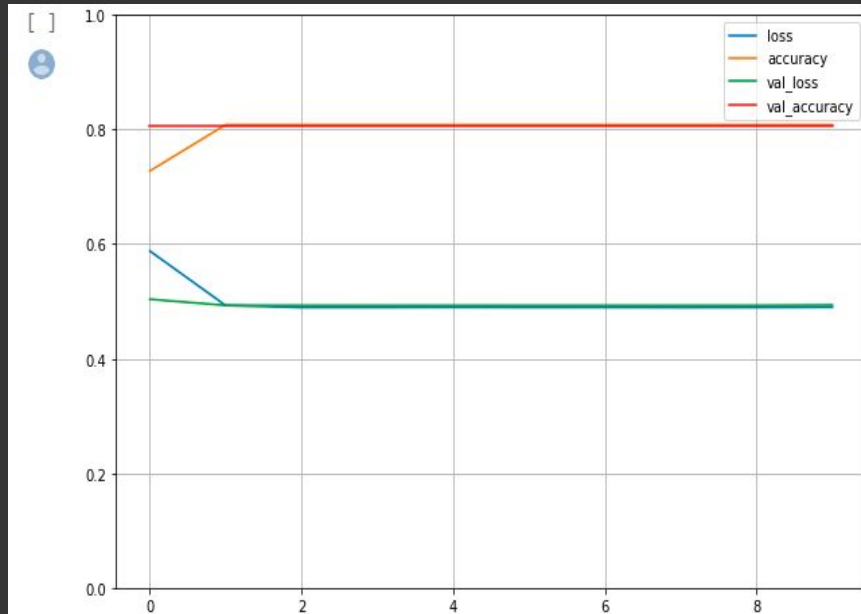


RESNet

- **Abnormal**

Loss : 0.5182

Accuracy : 0.7917

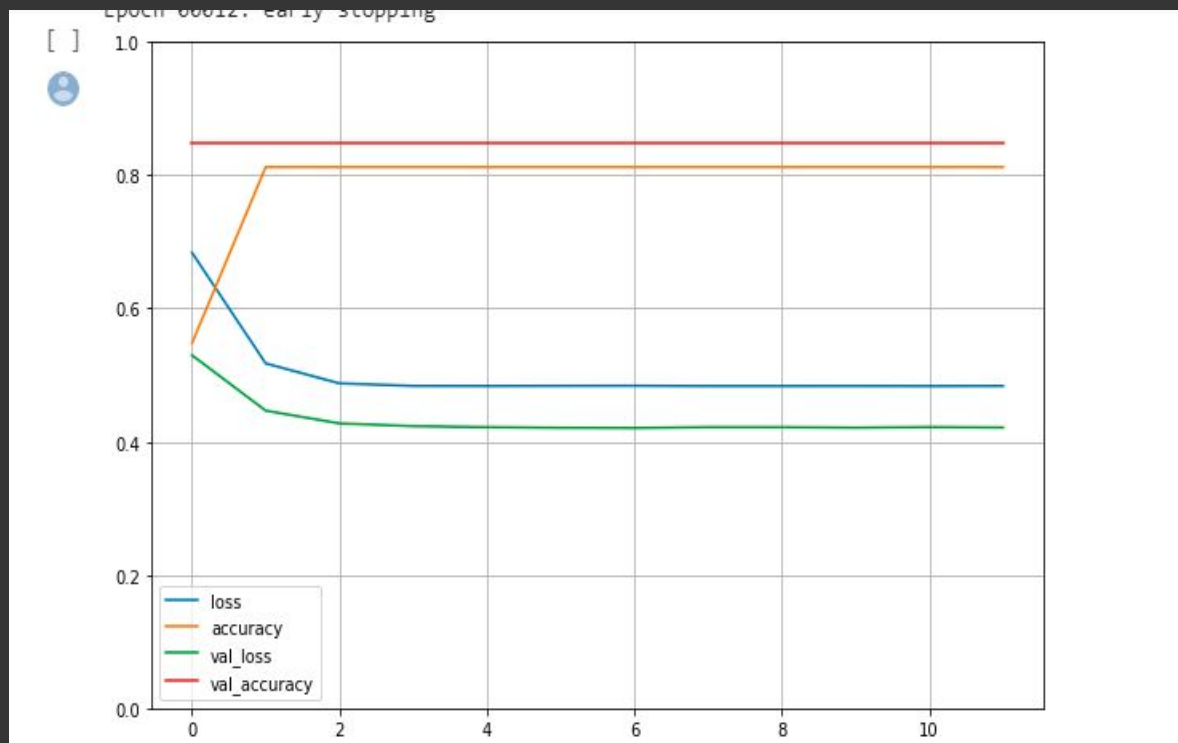


RESNet

- **ACL**

Loss : 0.7051

Accuracy : 0.5500



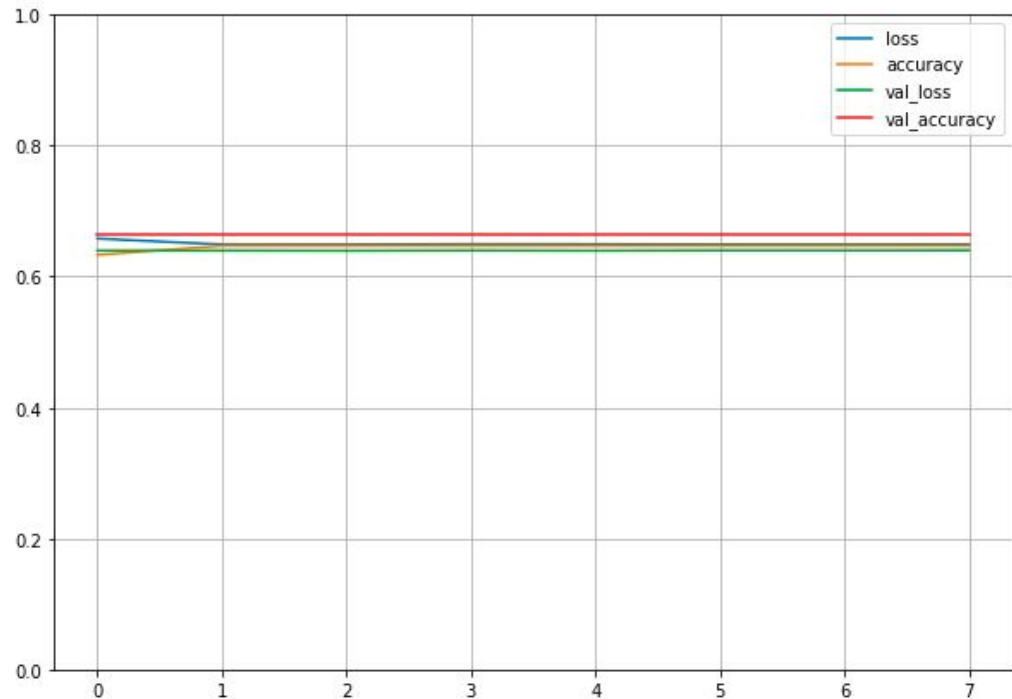
RESNet

- **Meniscal**

Loss : 0.7024

Accuracy : 0.5667

Epoch 00000: early stopping

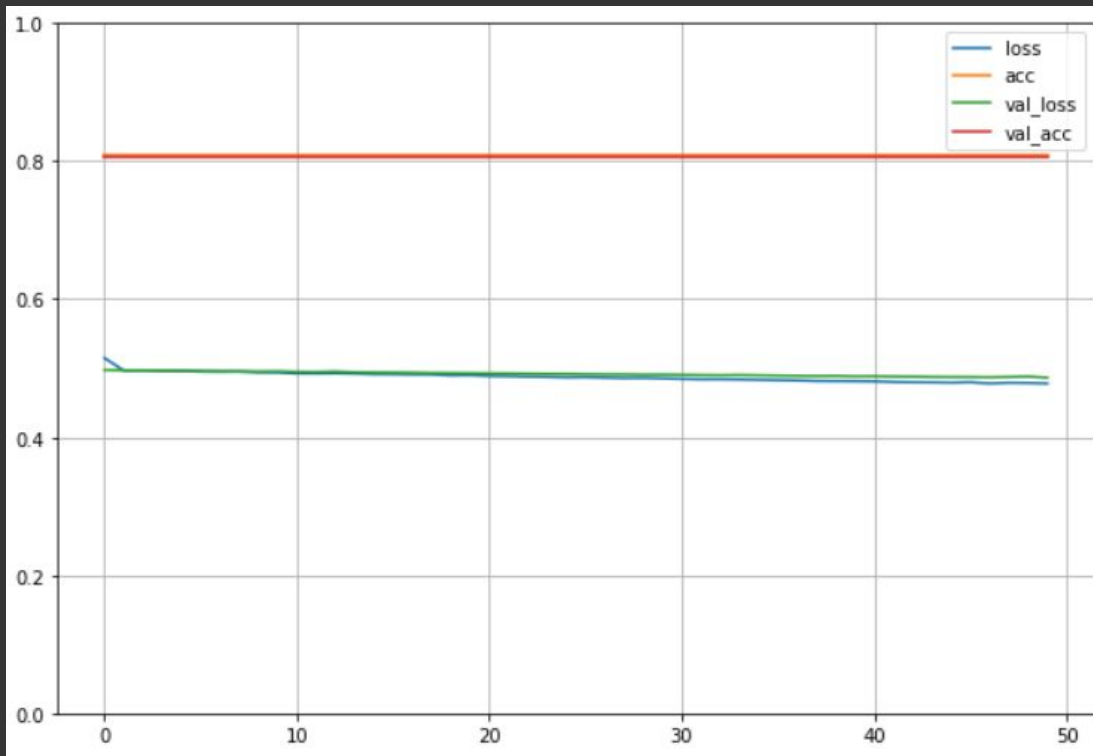


Inception V3

- **Abnormal**

Loss : 0.5167

Accuracy : 0.7917

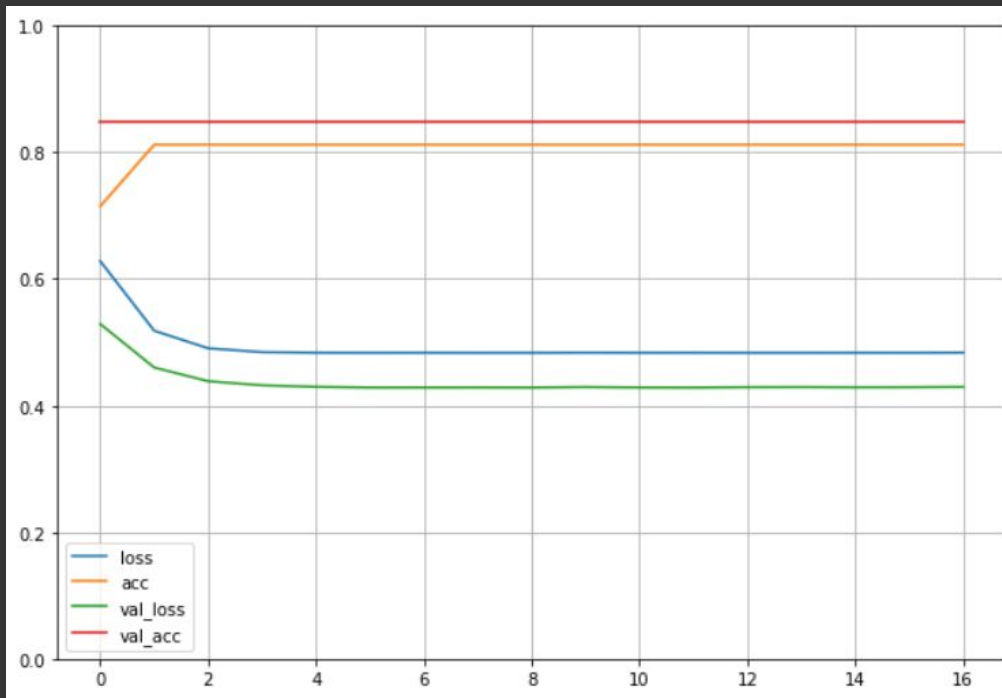


Inception V3

- **ACL**

Loss : 0.7071

Accuracy : 0.5500

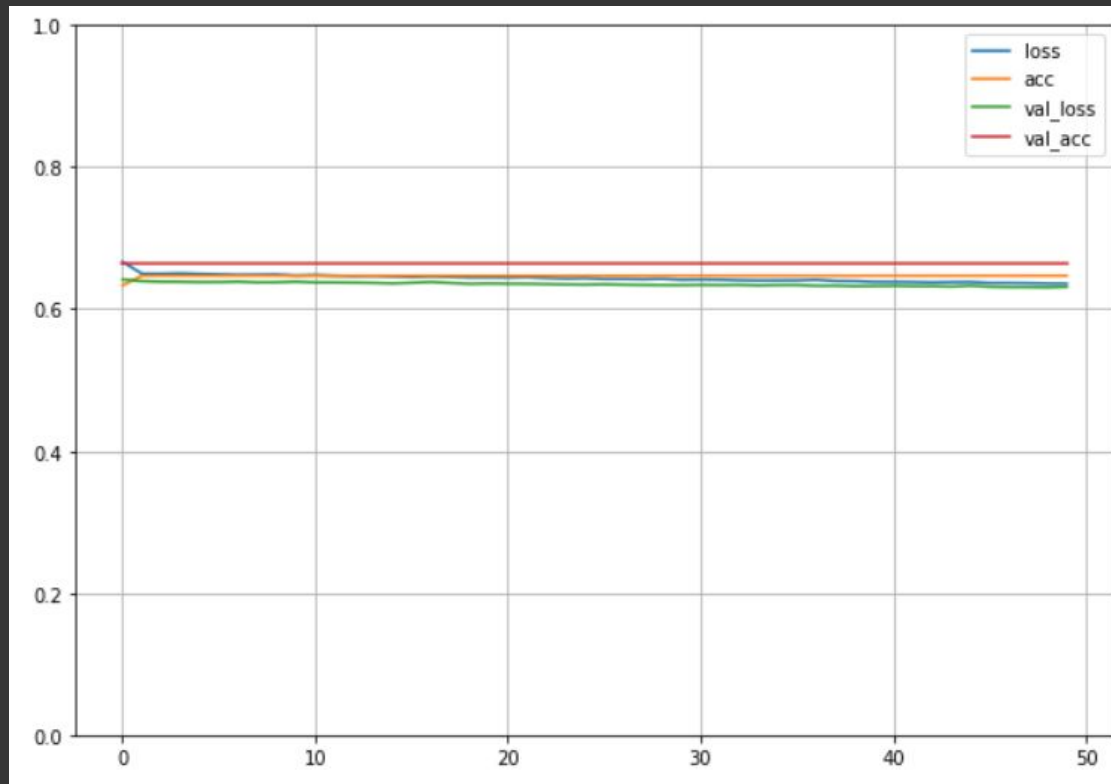


Inception V3

- **Meniscal**

Loss : 0.6912

Accuracy : 0.5667



models summary

	Anomaly	Accuracy	Loss
VGG	Abnormal	0.7917	0.5356
	Acl	0.5500	0.6903
	Meniscal	0.5667	0.6982
RESNet	Abnormal	0.7917	0.5182
	Acl	0.5500	0.7051
	Meniscal	0.5667	0.7024
Inception V3	Abnormal	0.7917	0.5167
	Acl	0.5500	0.7071
	Meniscal	0.5667	0.6912

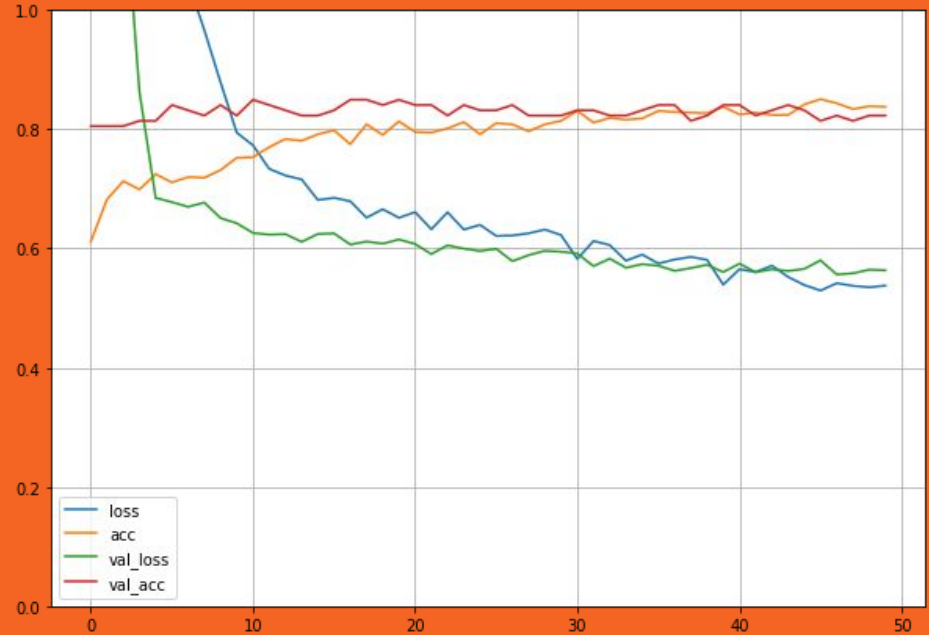
Transfer learning

We trained the three feature extractors (VGG , RESNet and Inception V3) by the models implemented in keras by training them on ImageNet which contains more than 14 million different images and then start training of classifiers and regressors by our Dataset.

Training Classifiers For VGG

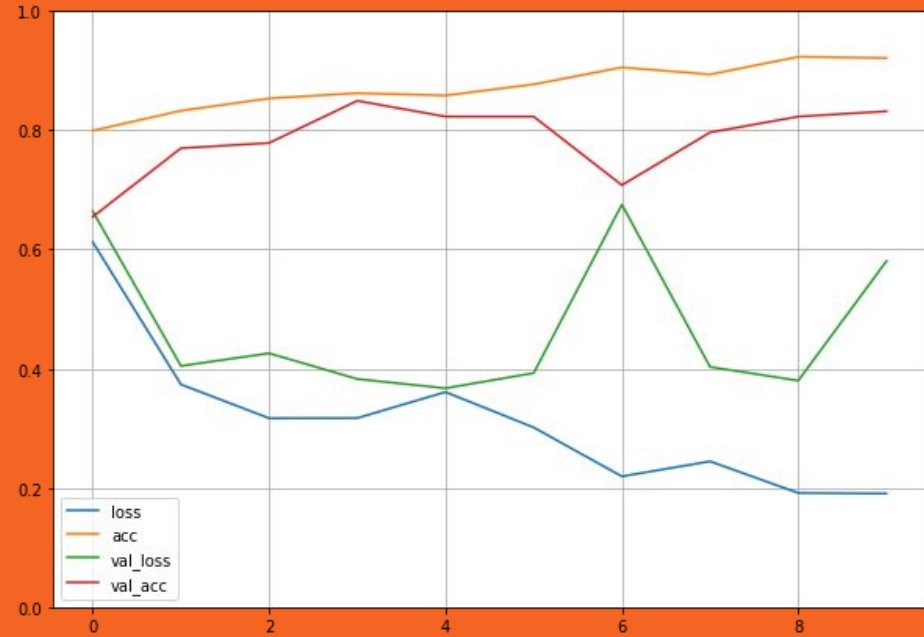
Classifier 1

Using less neurons in dense layers (128 neurons then 64 neurons with regularization and dropout layers)



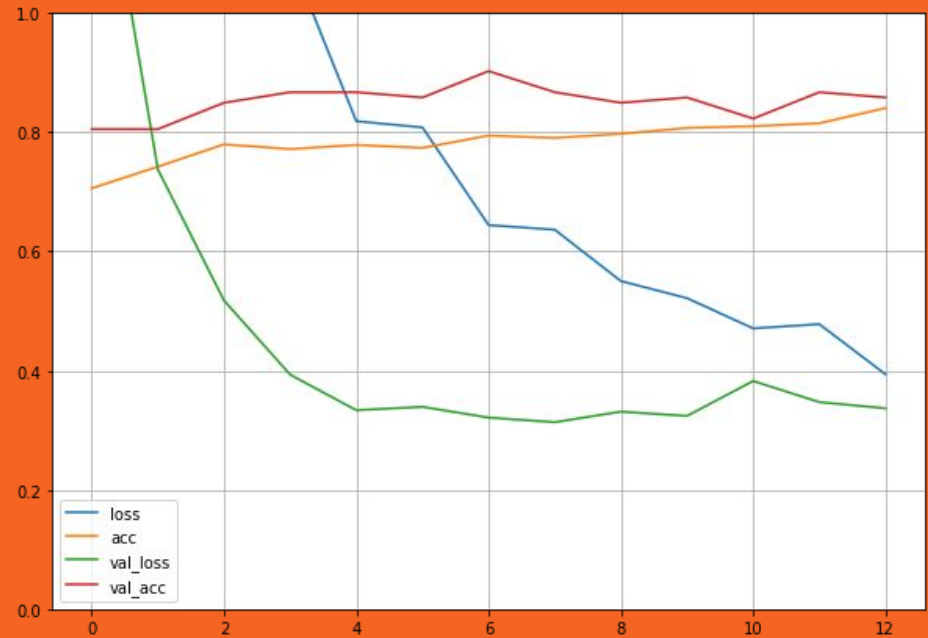
Classifier 2

Removing regularization and dropout layers and use more neurons in dense layers (1024 neurons then 512)



Classifier 3

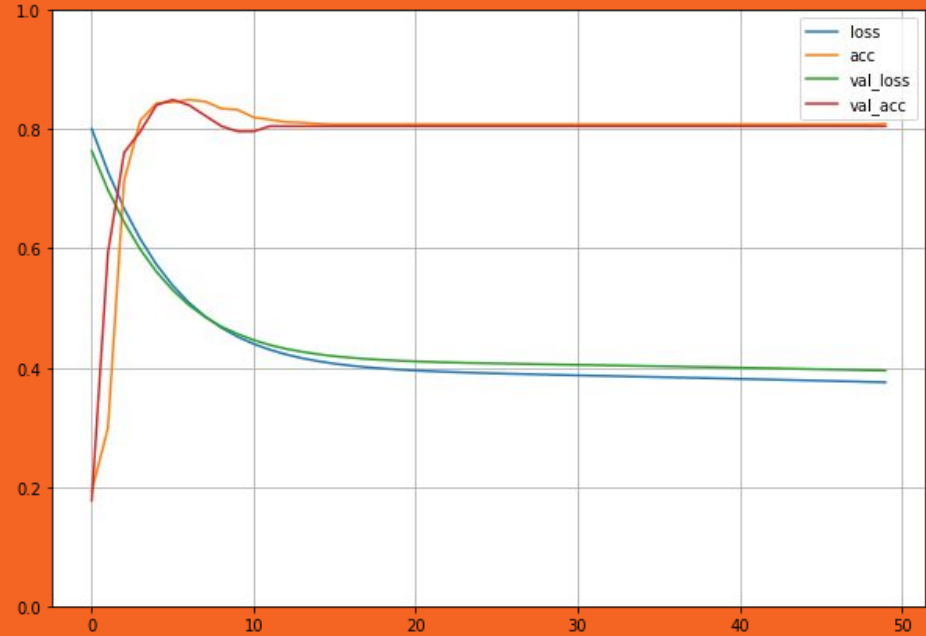
It is the final model for classifier by using more neurons in dense layers (1024 neurons then 512) and add dropout layers to overcome the overfitting



Training Regressors For VGG

Regressor 1

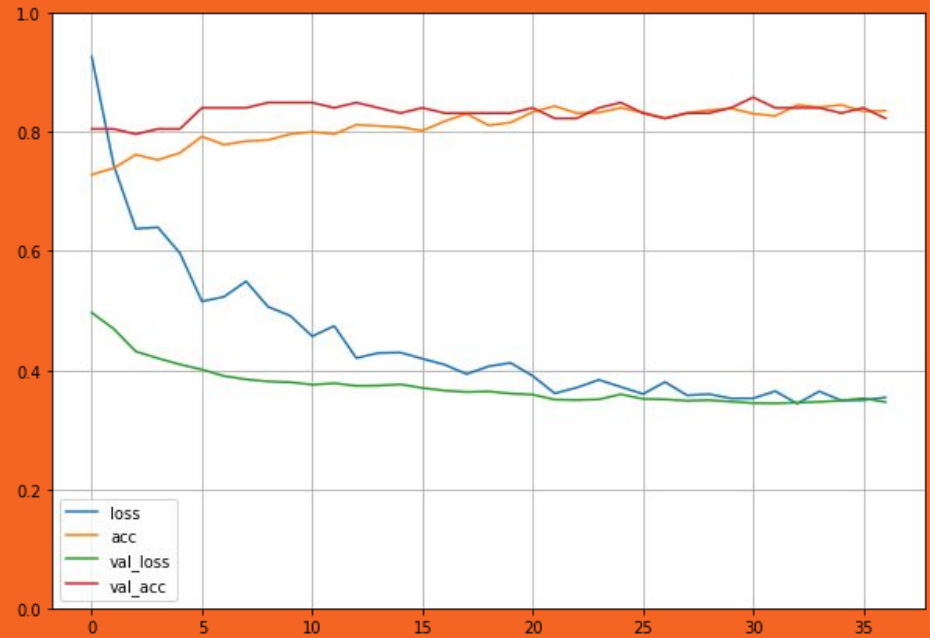
It is the final model for regressor using the adam optimizer with learning rate = 0.01



Training Classifiers For RESNet

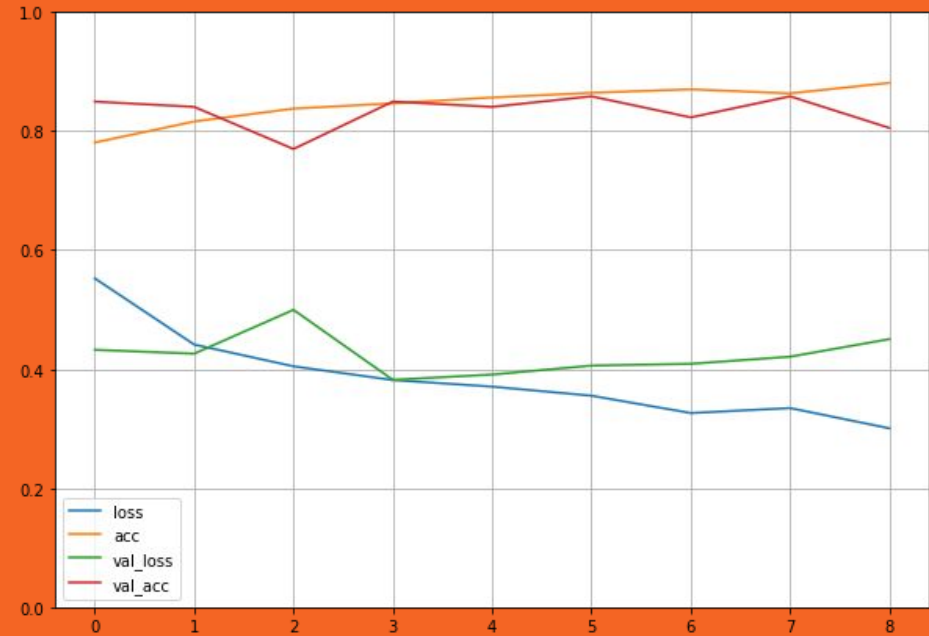
Classifier 1

Using less neurons in dense layers (512 neurons then 256 neurons with dropout layers)



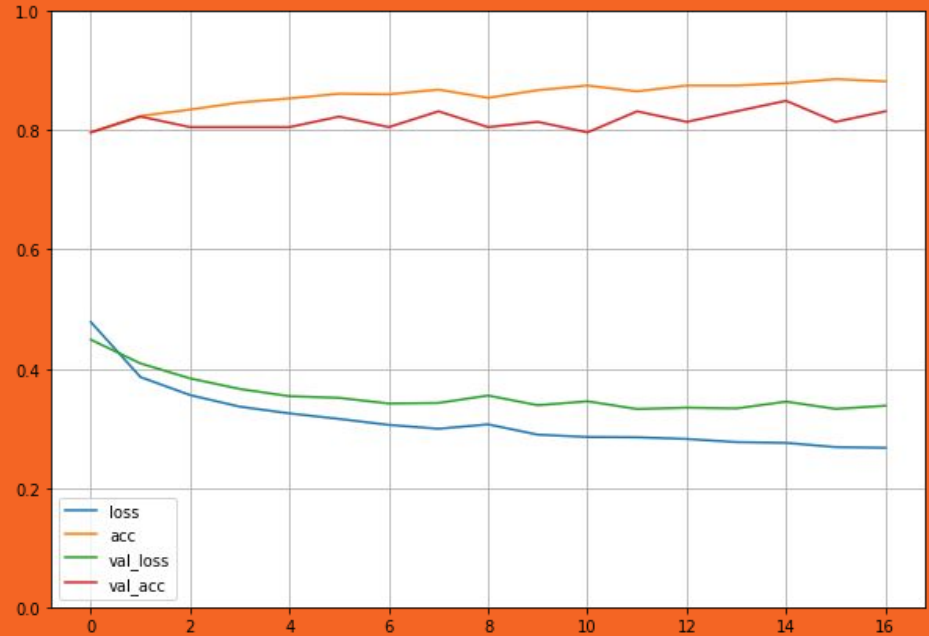
Classifier 2

Using learning rate = 10^{-4} and this test is for Axial ACL classifier not Axial Abnormal like the other ones



Classifier 3

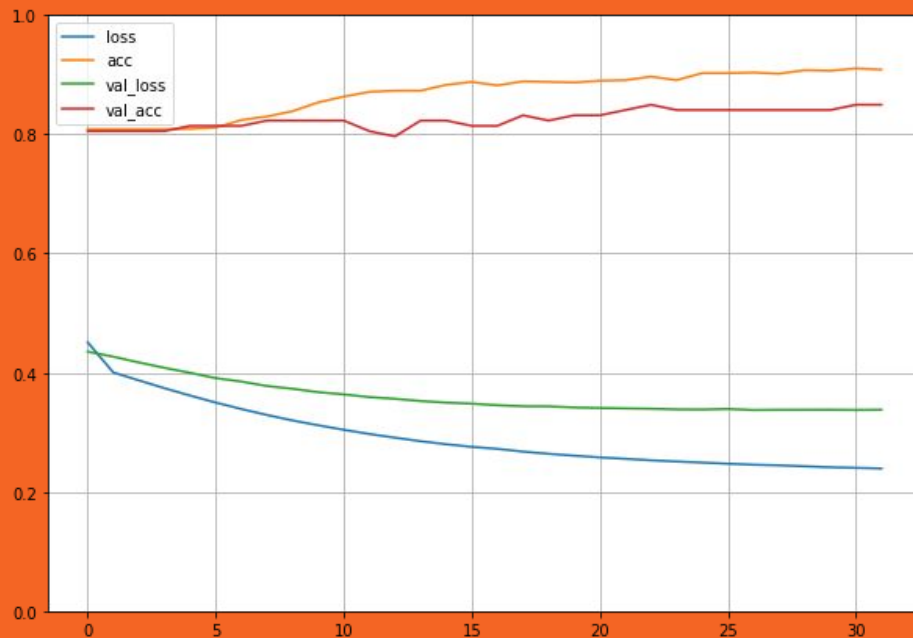
It is the final model for the classifier by using learning rate = 10^{-5} for adam optimizer with (512 neurons in first dense and 256 in the next dense layer)



Training Regressors For RESNet

Regressor 1

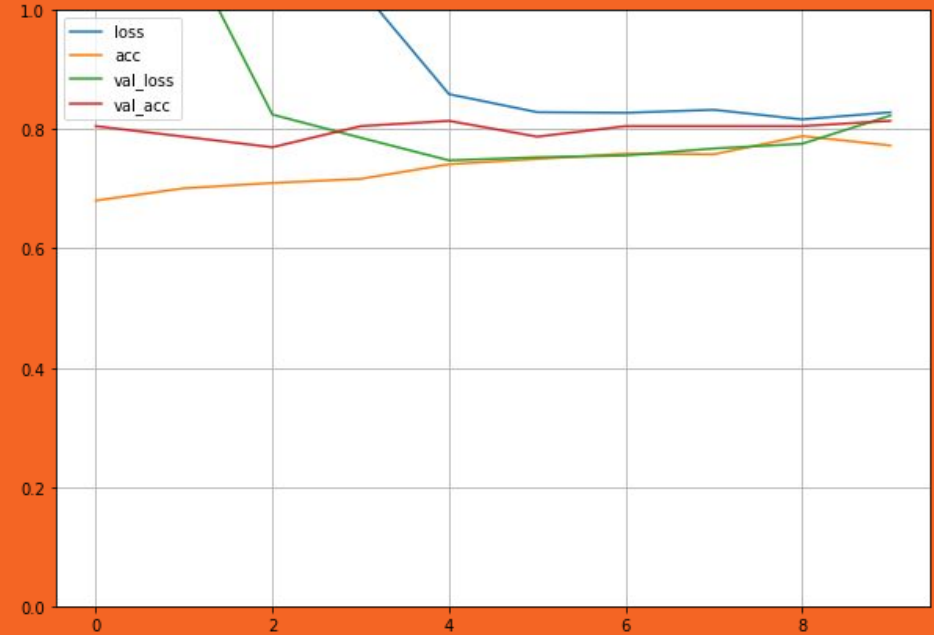
It the final model for regressor using the learning rate = 0.01 by adam optimizer



Training Classifiers For Inception V3

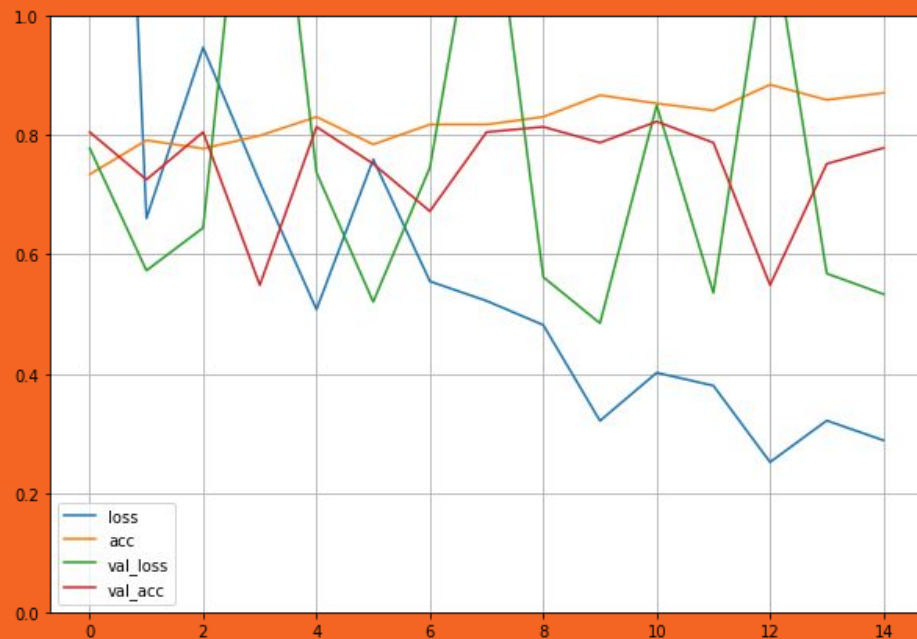
Classifier 1

Using less neurons in dense layers (128 neurons then 64 neurons as MRNet paper with regularization and dropout layers)



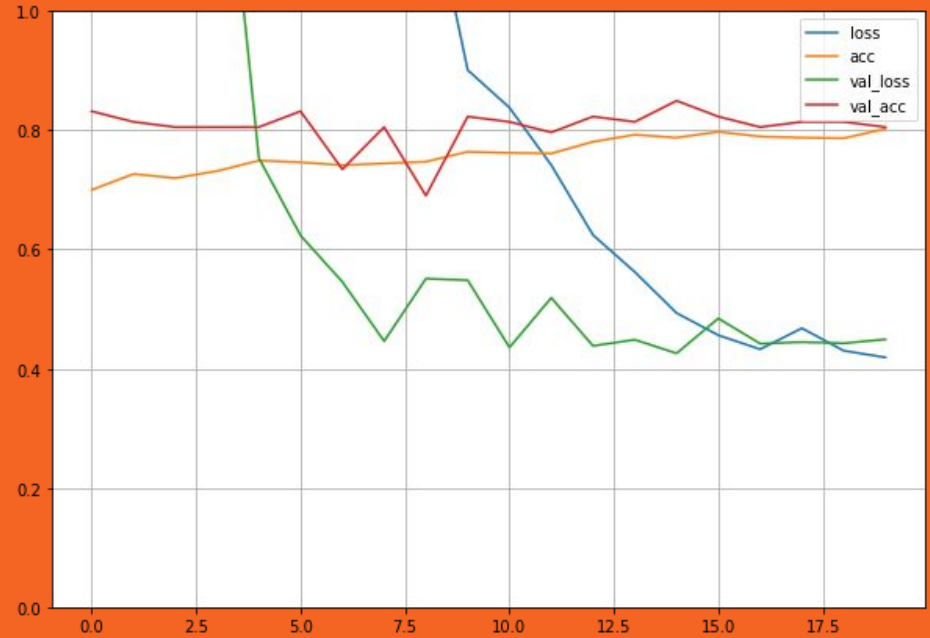
Classifier 2

Removing regularization and dropout layers and use more neurons in dense layers (1024 neurons then 512)



Classifier 3

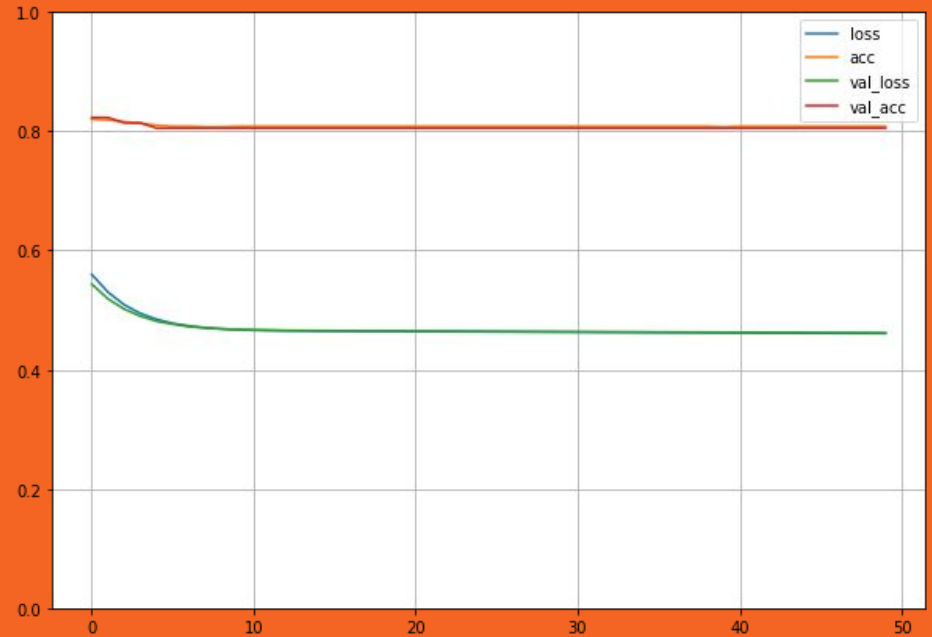
It is the final model for classifier by using more neurons in dense layers (1024 neurons then 512) and add dropout layers to overcome the overfitting



Training Regressors For Inception V3

Regressor 1

It the final model for regressor using the learning rate = 0.01 by adam optimizer



Transfer Learning Statistics



Regressors

We judge the model by two methods

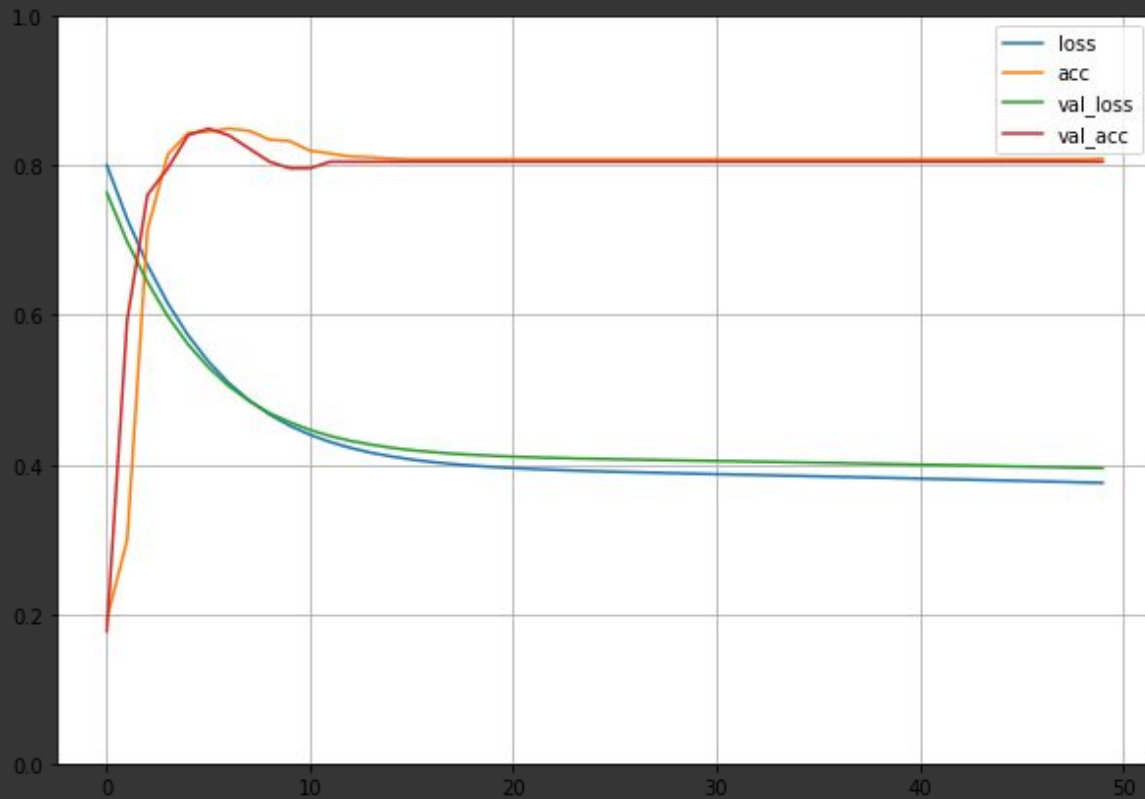
- Loss
- Accuracy

VGG

- **Abnormal**

Loss : 0.5338

Accuracy : 0.8083

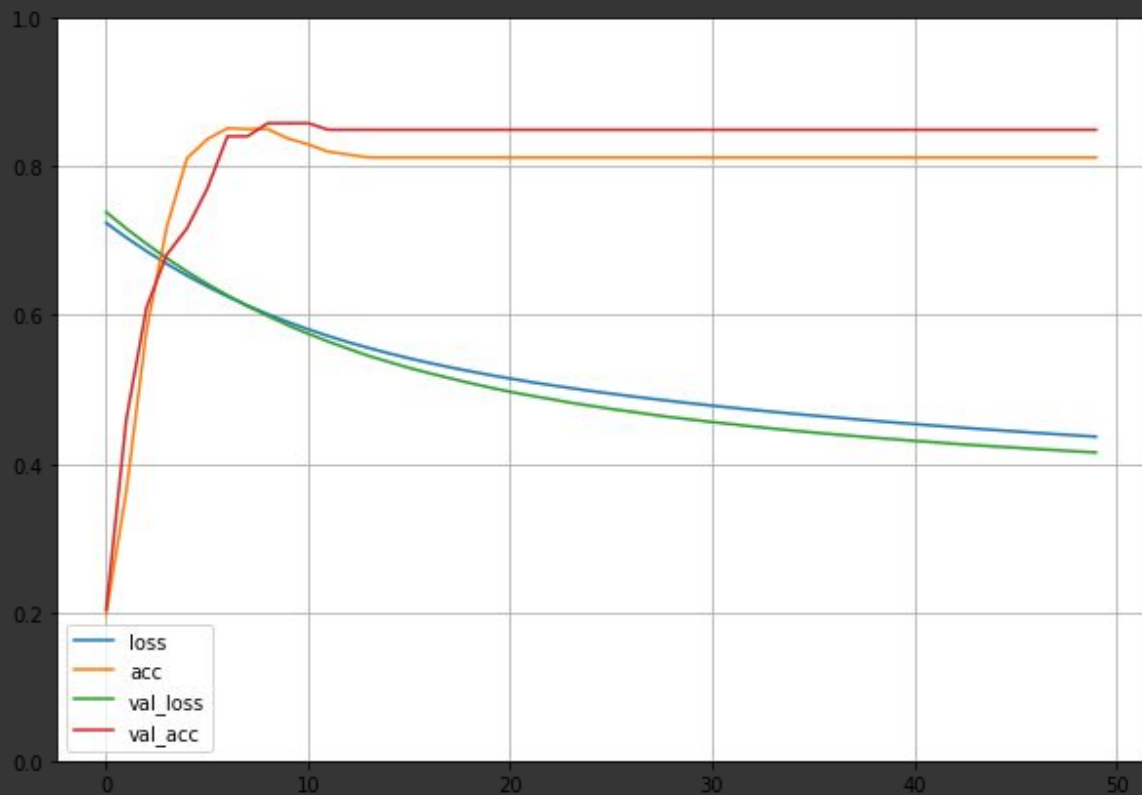


VGG

- **ACL**

Loss : 0.659

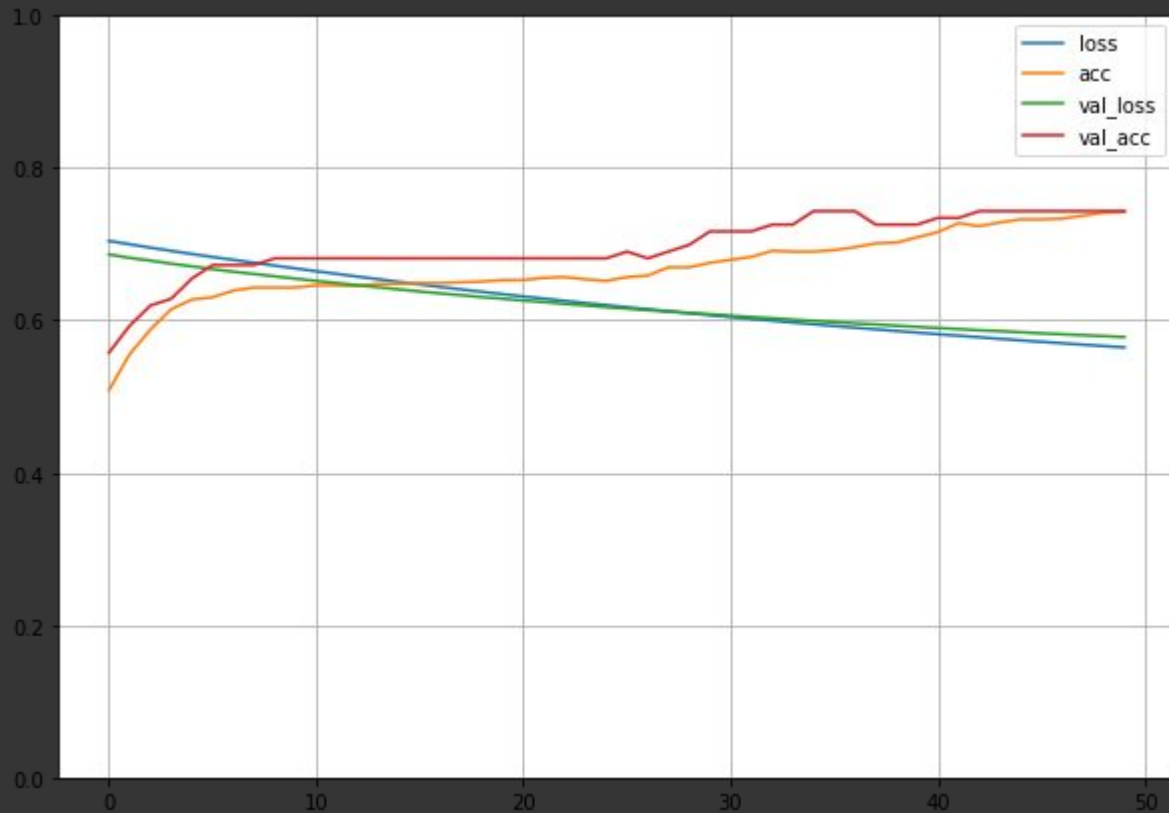
Accuracy : 0.625



VGG

- **Meniscal**

Loss : 0.6374
Accuracy : 0.5917

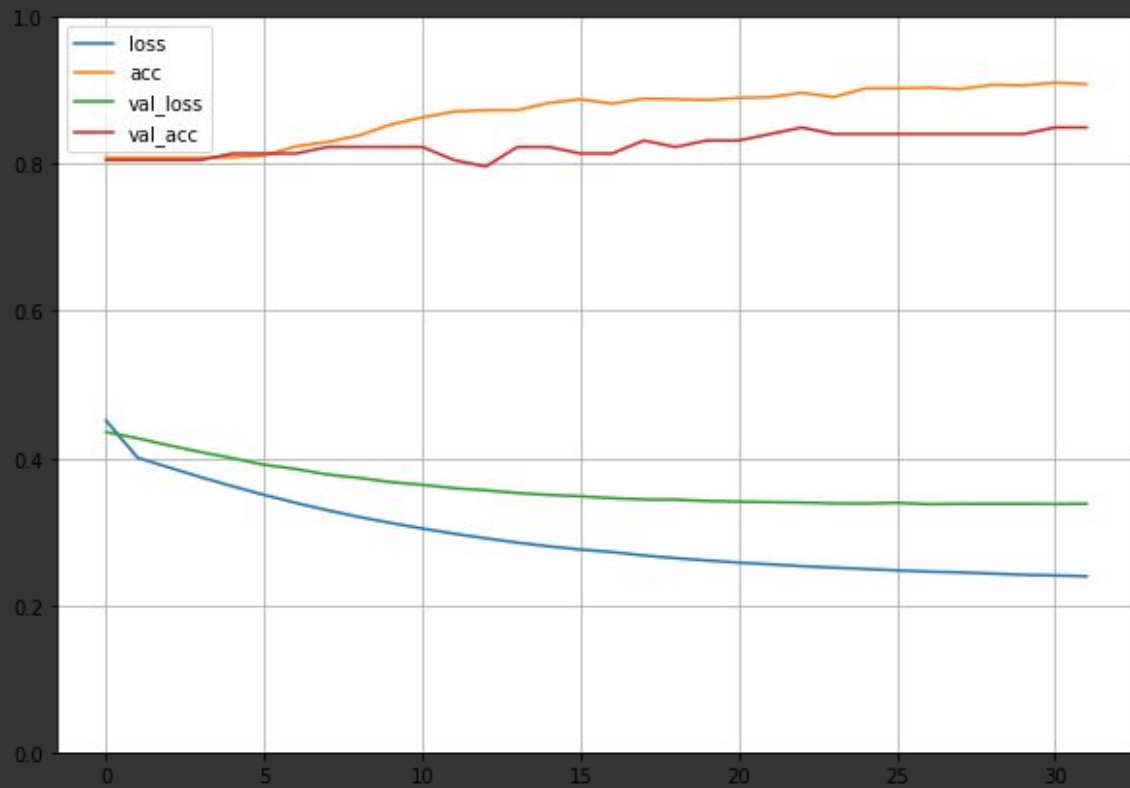


RESNet

- **Abnormal**

Loss : 0.3940

Accuracy : 0.8333

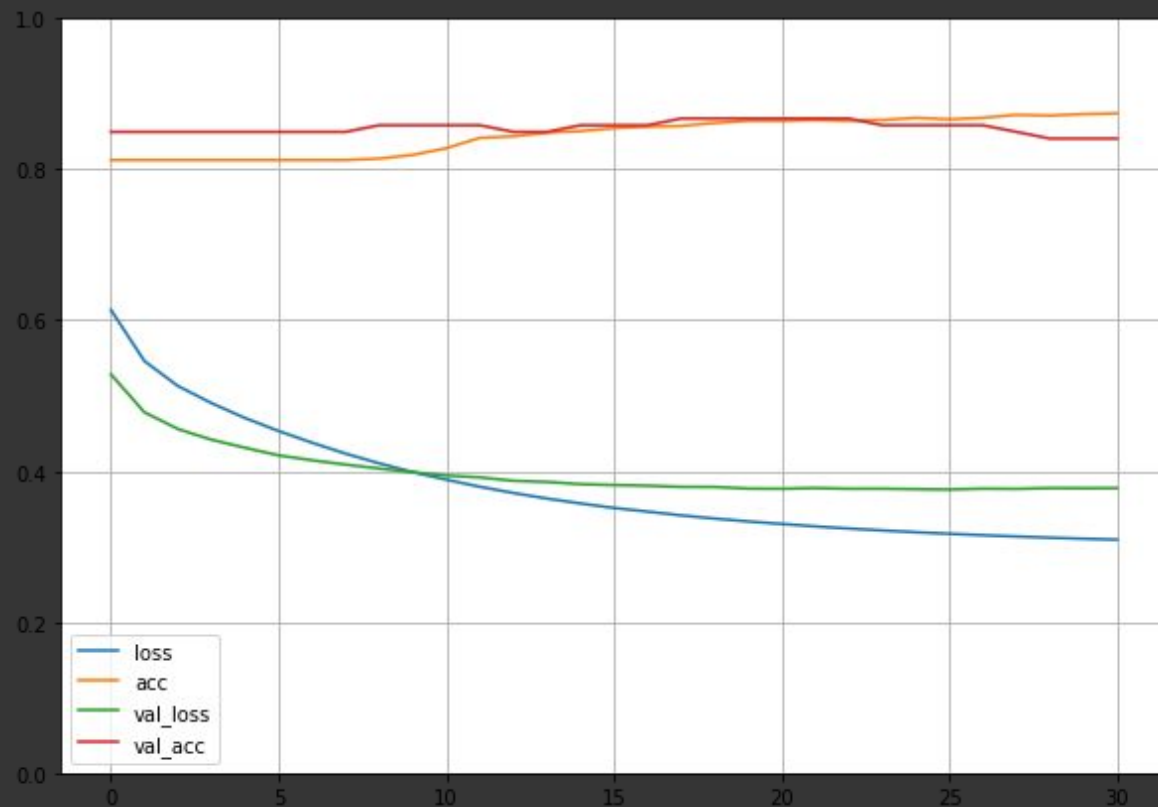


RESNet

- **ACL**

Loss : 0.6274

Accuracy : 0.6583

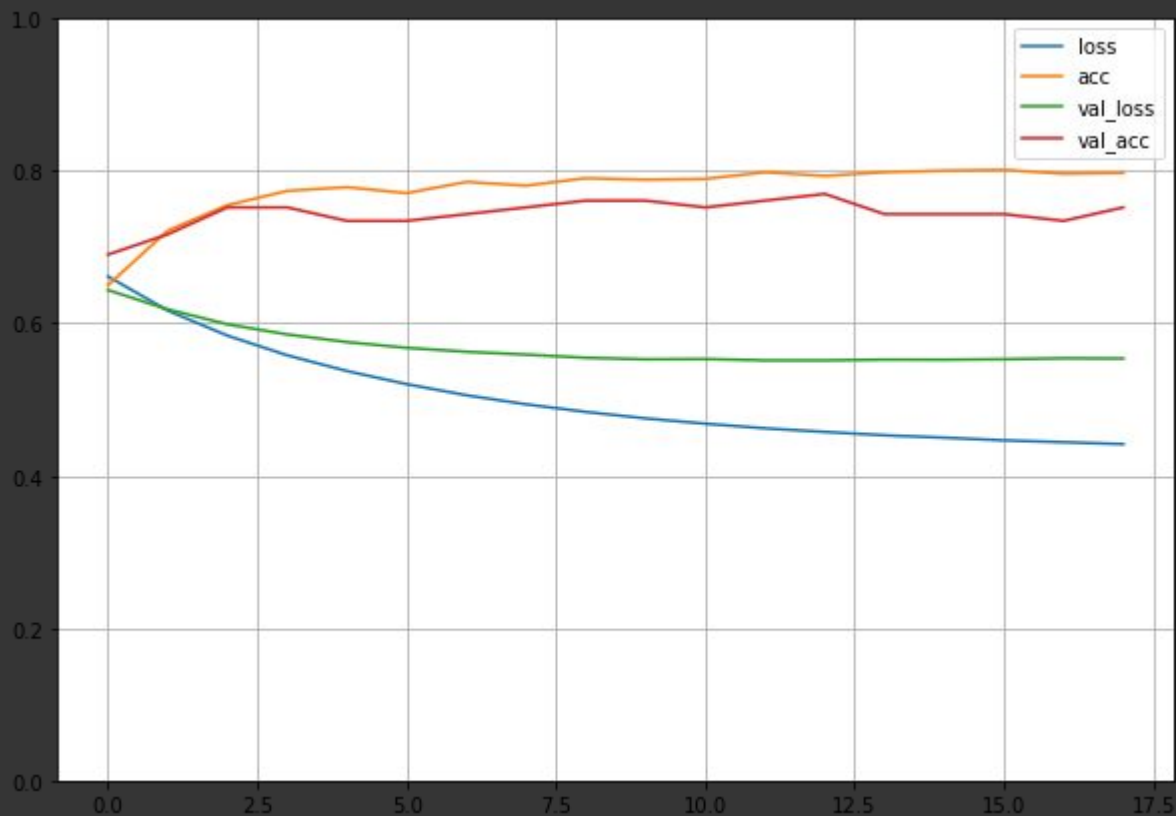


RESNet

- **Meniscal**

Loss : 0.5947

Accuracy : 0.7083

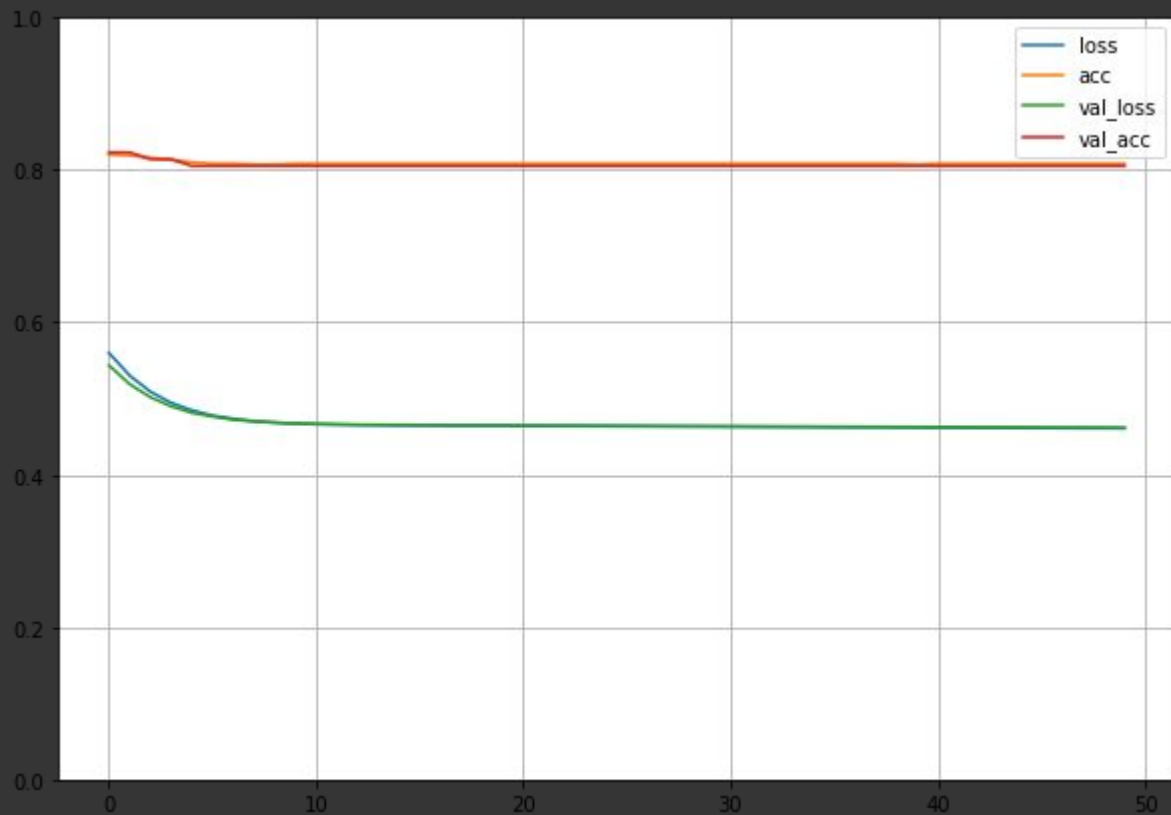


Inception V3

- **Abnormal**

Loss : 0.5504

Accuracy : 0.7917

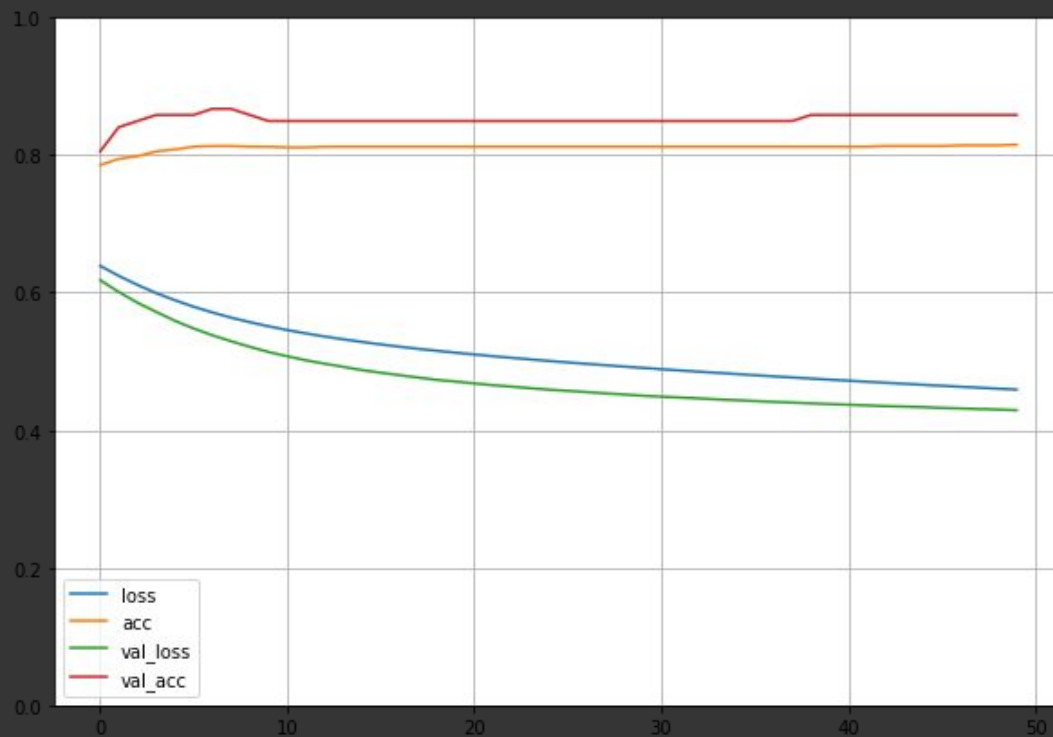


Inception V3

- **ACL**

Loss : 0.7141

Accuracy : 0.5417

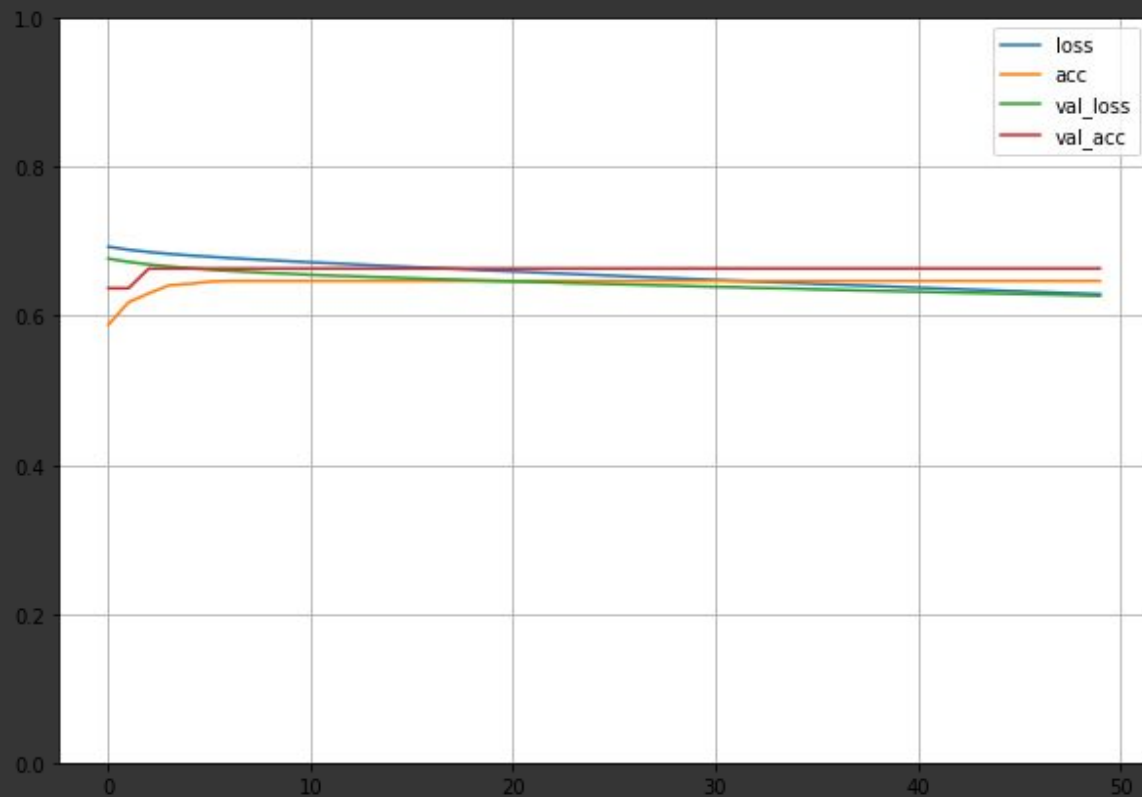


Inception V3

- **Meniscal**

Loss : 0.7000

Accuracy : 0.5667



Transfer learning models summary

	Anomaly	Accuracy	Loss
VGG	Abnormal	0.8083	0.5338
	Acl	0.6250	0.6590
	Meniscal	0.5917	0.6374
RESNet	Abnormal	0.8333	0.3940
	Acl	0.6583	0.6274
	Meniscal	0.7083	0.5947
Inception V3	Abnormal	0.7917	0.5504
	Acl	0.5417	0.7141
	Meniscal	0.5667	0.7000

Contribution



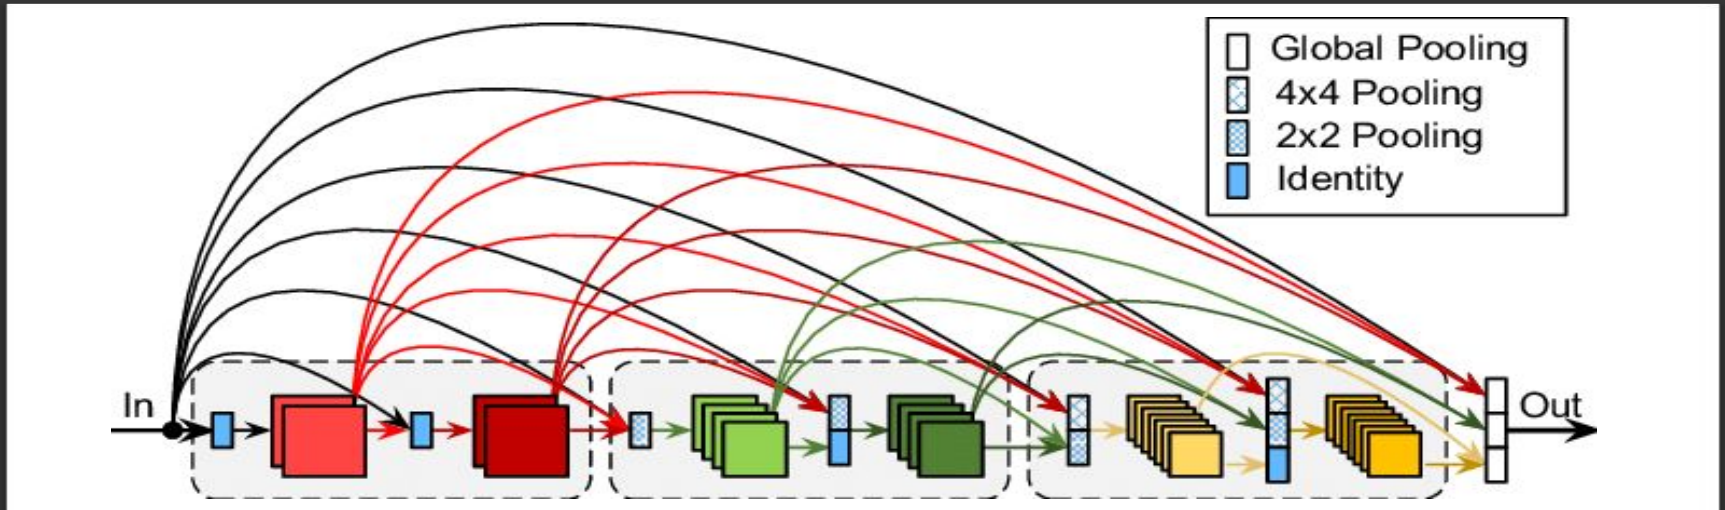
The first paper

“ Deep Learning for Musculoskeletal Image Analysis” suggests two approaches

- Train DenseNet model as feature extractor By transfer learning
- Use large amount of imaging data (sooo... challenging !!!)

DenseNet

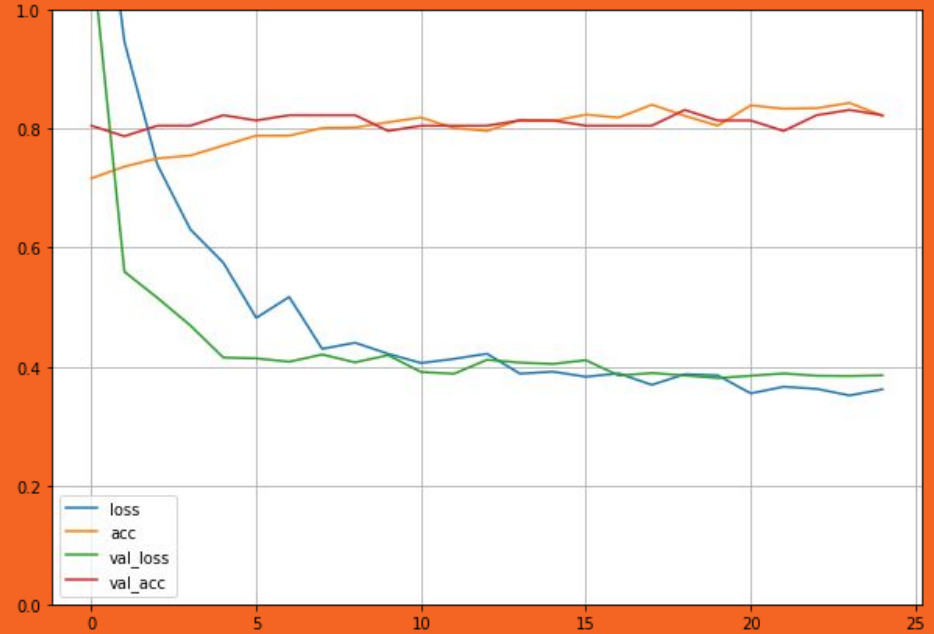
- It enhances the RESNet model by connecting output of a layer to all subsequent layers. This architecture helps in improving in the flow of information and gradient in the network.



Training Classifiers For DenseNet

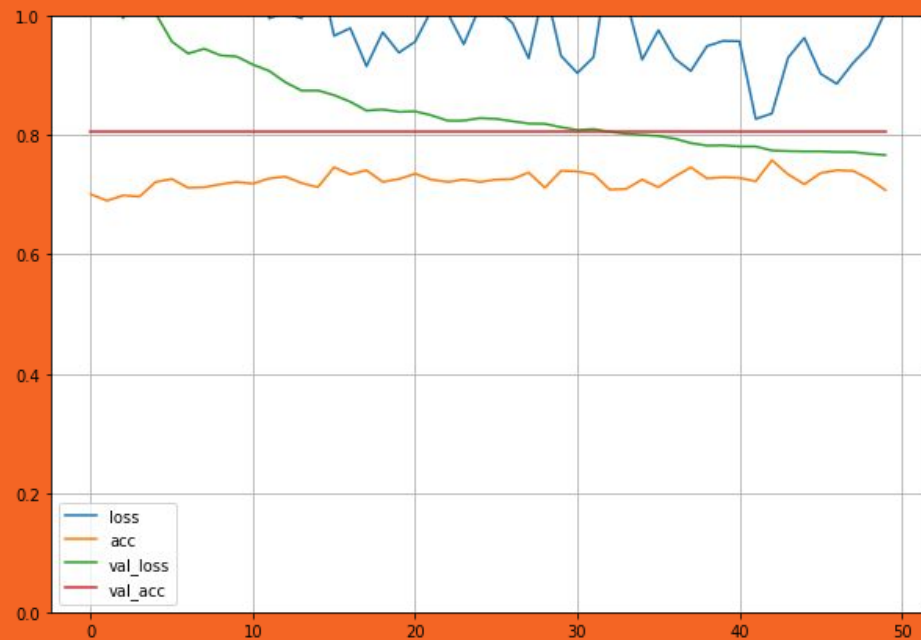
Classifier 1

Using less neurons in dense layers (512 neurons then 256 neurons with dropout layers)



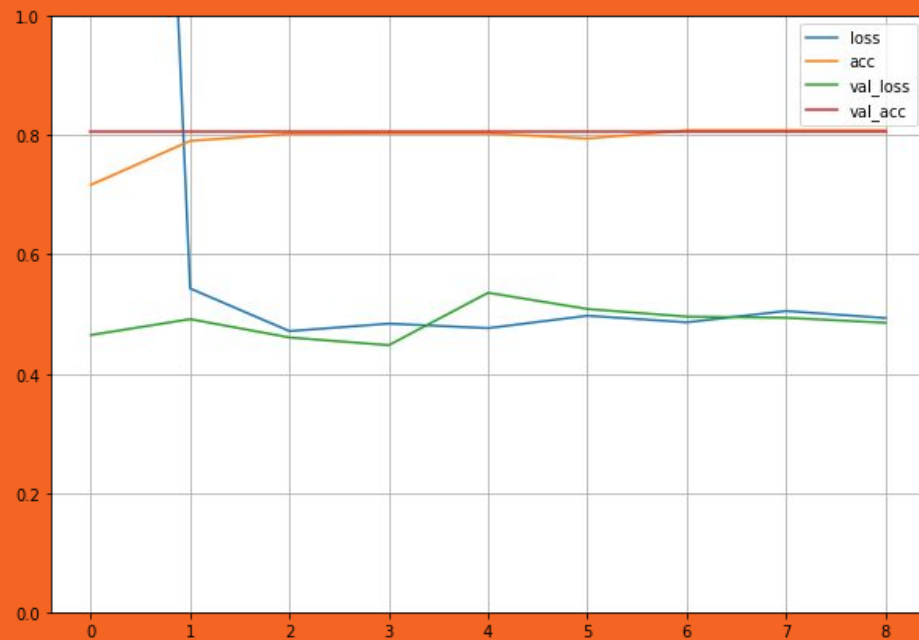
Classifier 2

Using learning rate = 10^{-4} and decay rate = 0.1



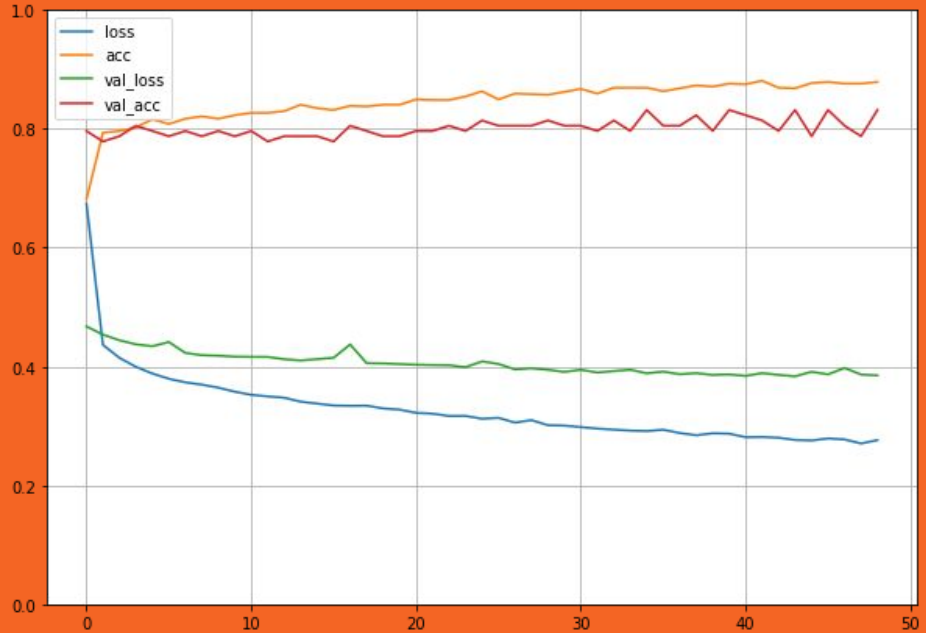
Classifier 3

Increase learning rate to 0.01



Classifier 4

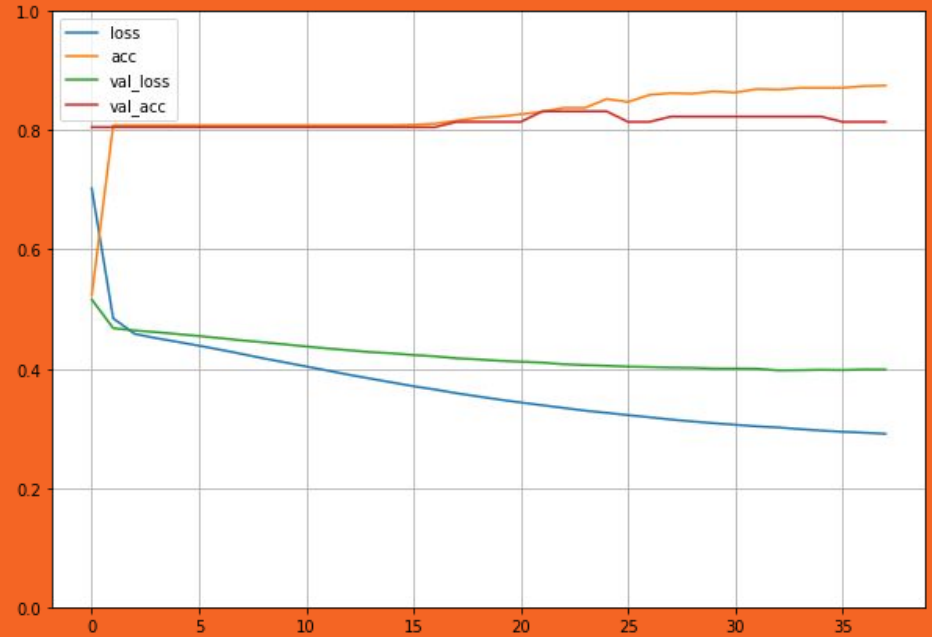
It is the final model for the classifier by using learning rate = 10^{-5} for adam optimizer with (512 neurons in first dense and 256 in the next dense layer)



Training Regressors For DenseNet

Regressor 1

It the final model for regressor using the learning rate = 0.01 by adam optimizer



DenseNet Model Statistics



Regressors

We judge the model by two methods

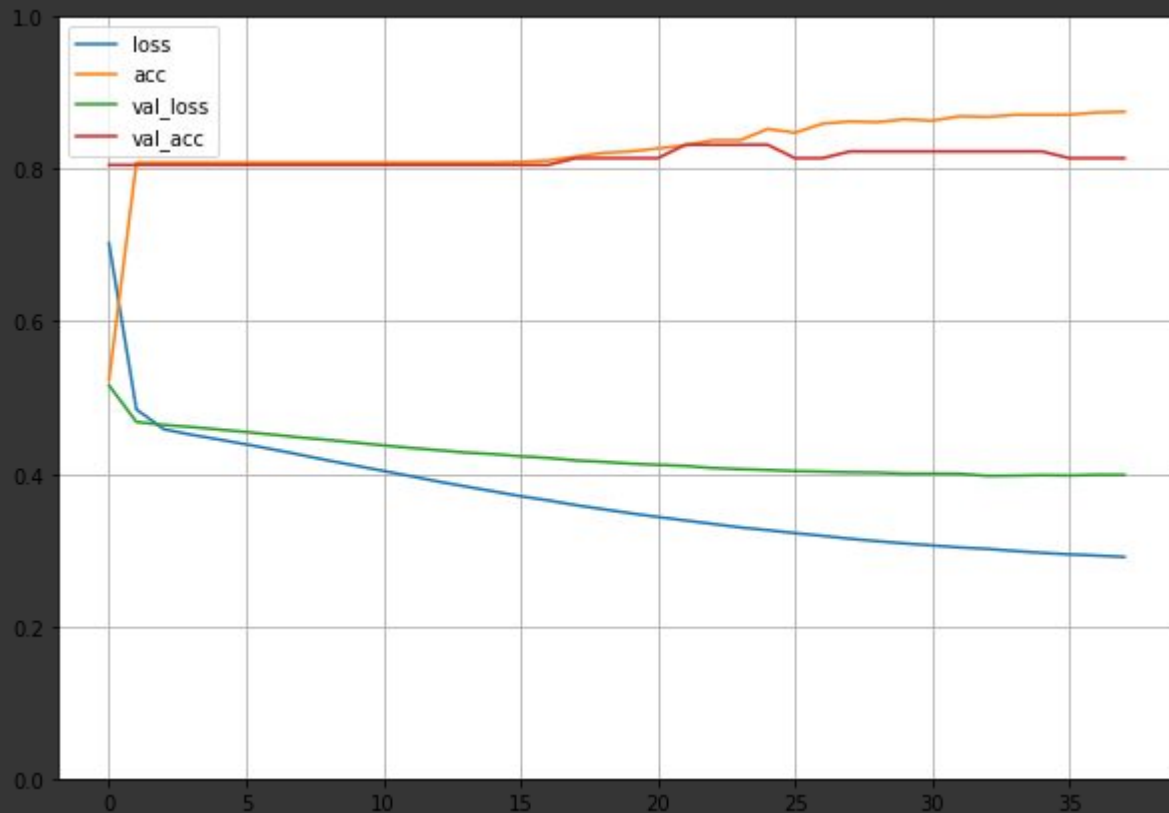
- Loss
- Accuracy

DenseNet

- **Abnormal**

Loss : 0.4418

Accuracy : 0.8167

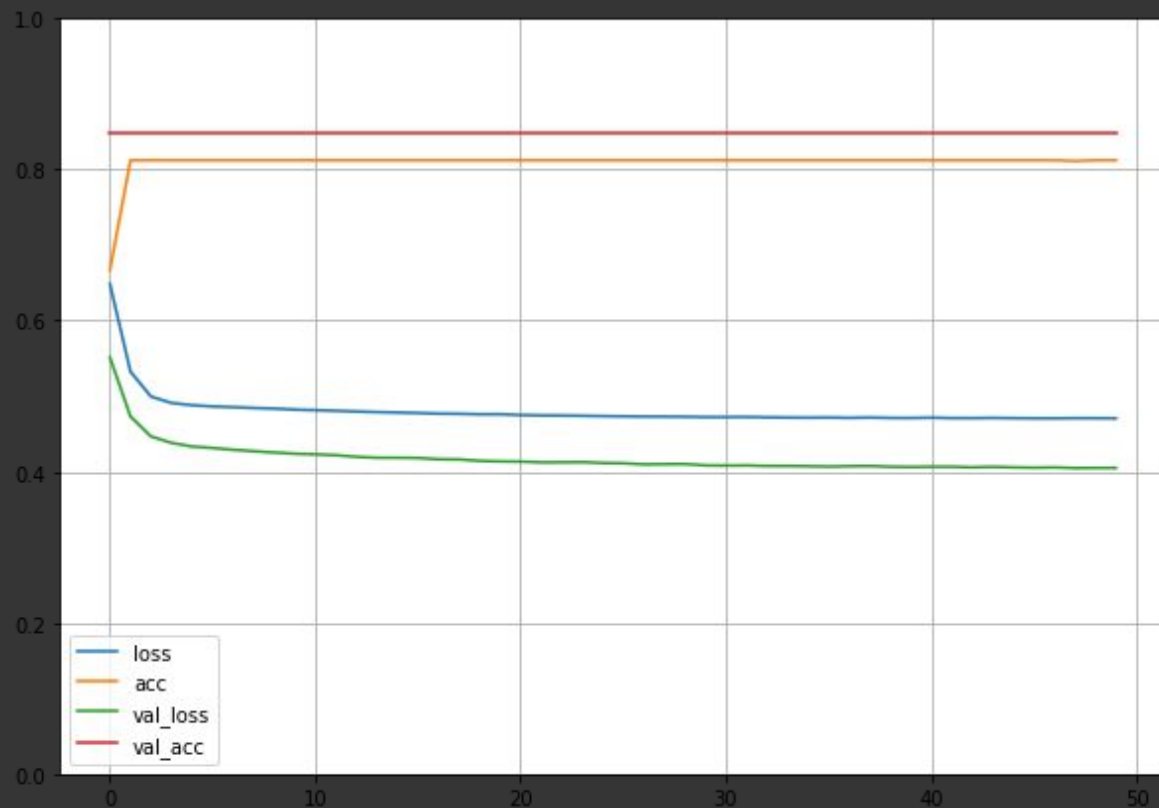


DenseNet

- **ACL**

Loss : 0.6972

Accuracy : 0.5500

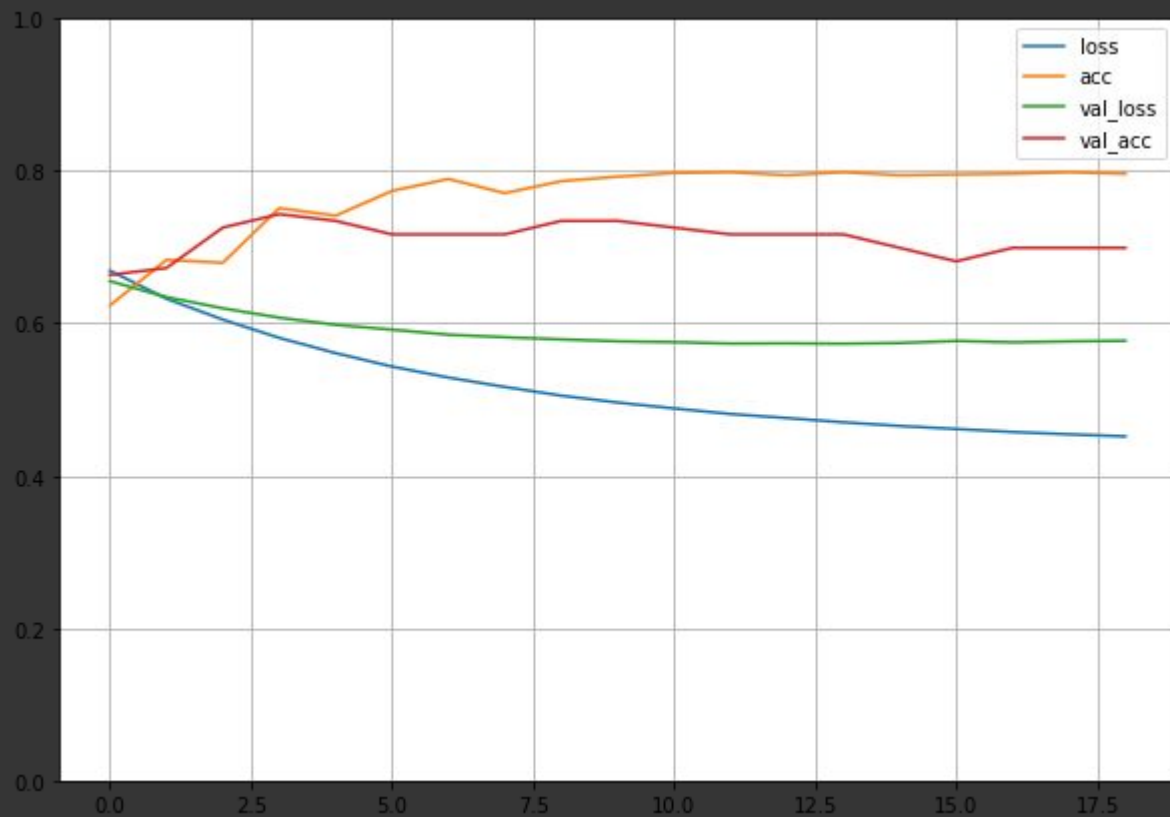


DenseNet

- **Meniscal**

Loss : 0.6453

Accuracy : 0.6417



DenseNet model summary

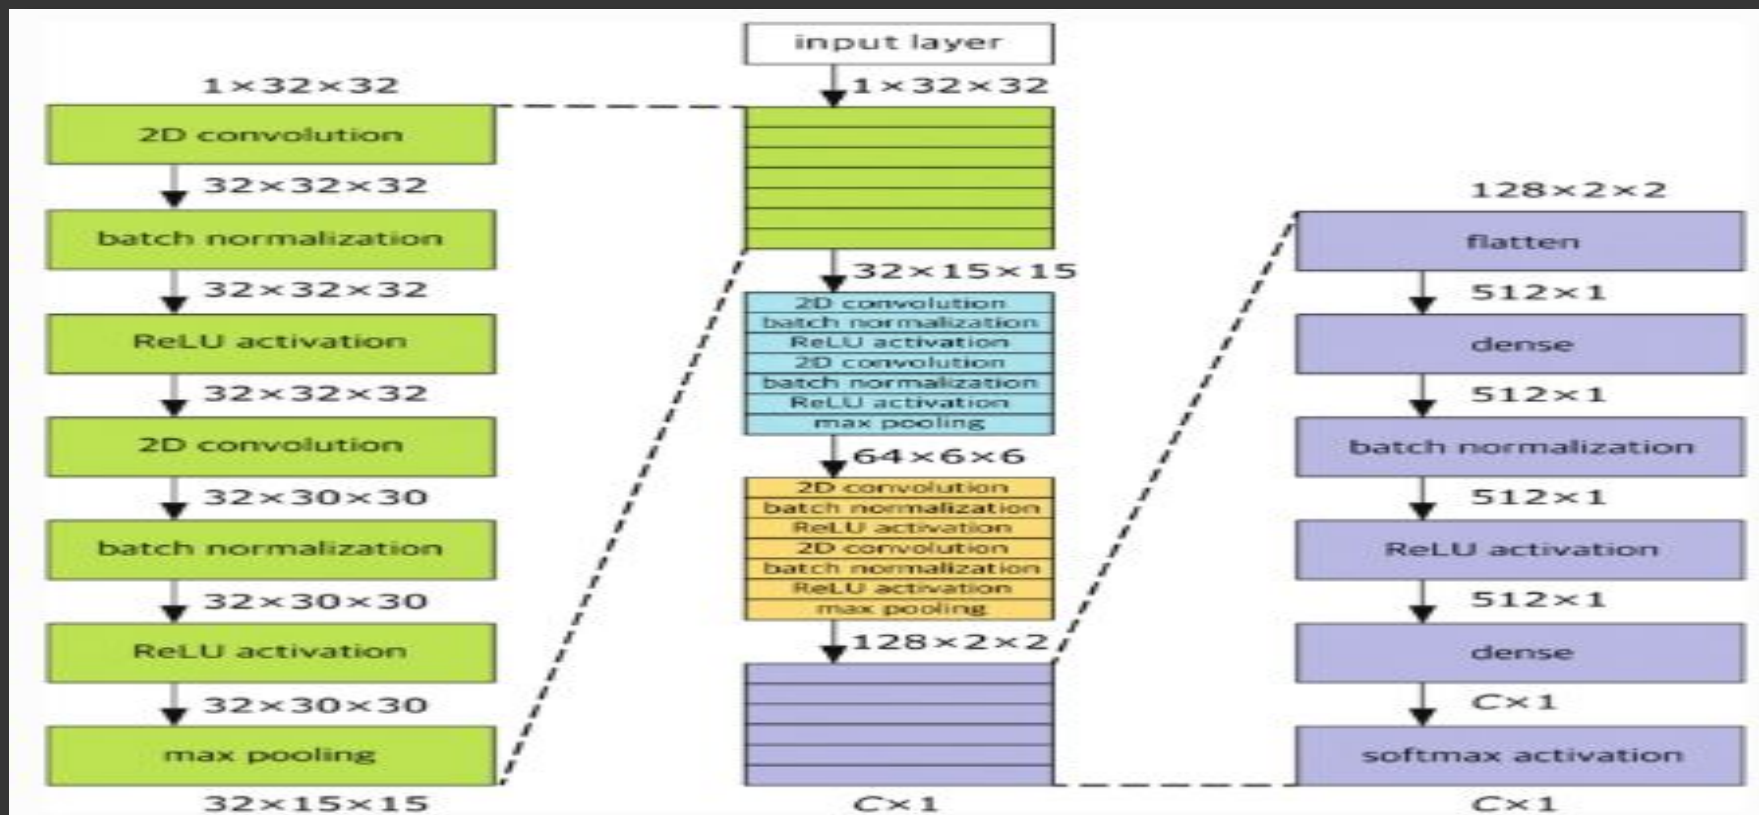
	Accuracy	Loss
Abnormal	0.8167	0.4418
ACL	0.5500	0.6972
Meniscal	0.6417	0.6453



The second paper

**Using Deep Learning Algorithms to
Automatically Identify the Brain MRI
Contrast: Implications for Managing Large
Databases**

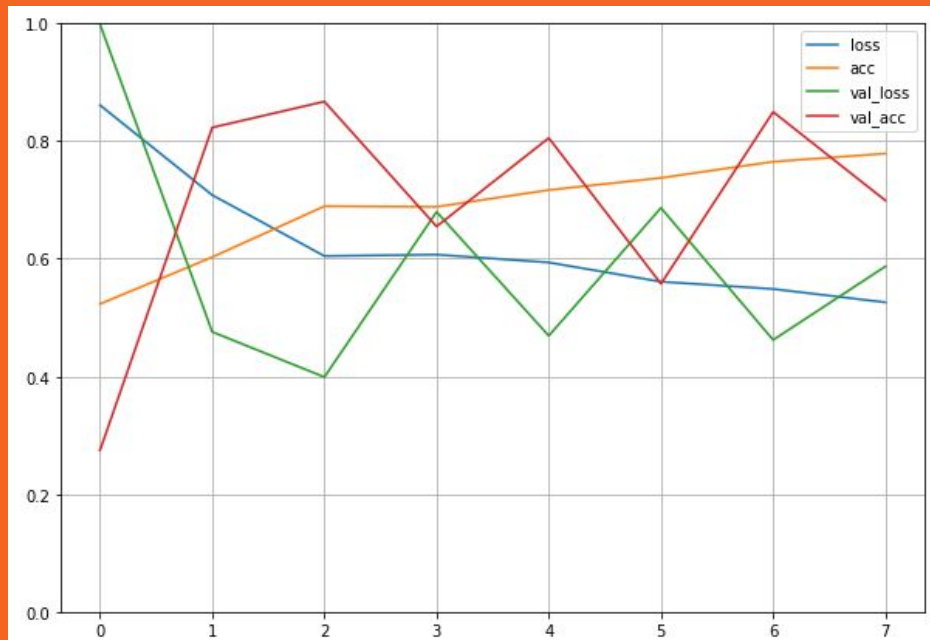
Contrast



Training Extractors For Contrast

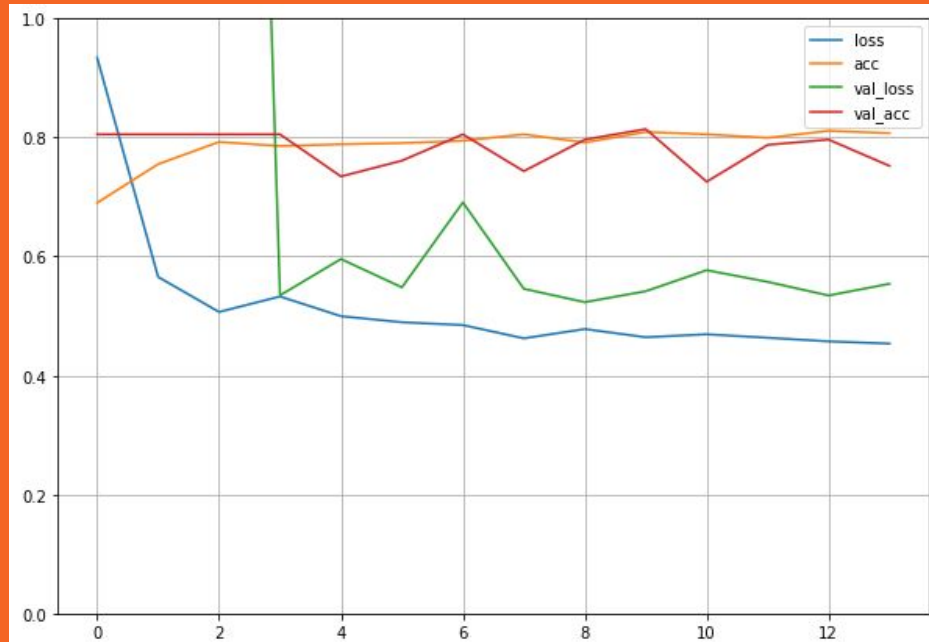
Extractor 1

Without using dropout layer and using adam optimizer there was overfitting problem



Extractor 2

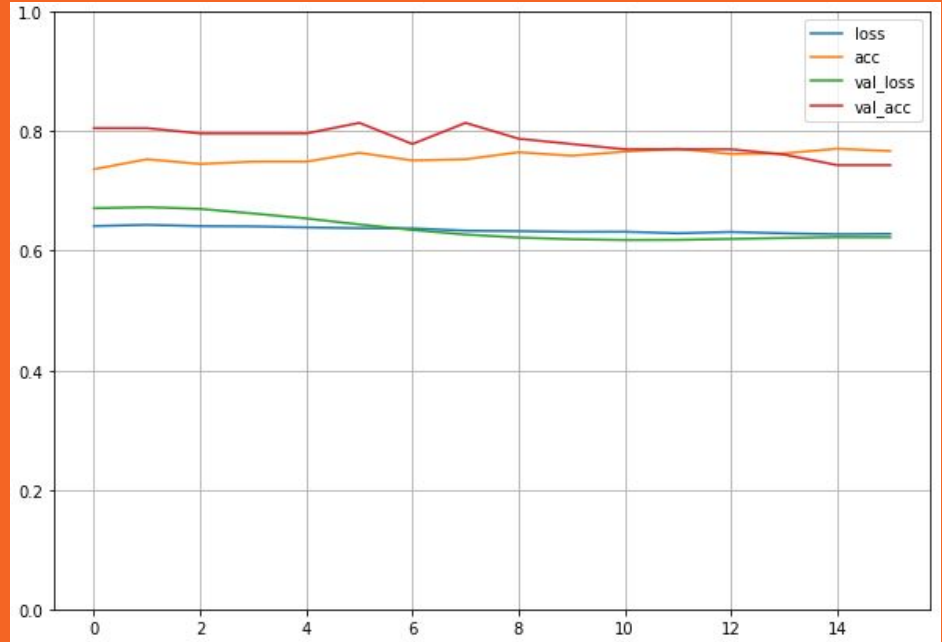
Using dropout layer reduces the overfitting problem also optimizer changed to nadam



Training Classifiers For Contrast

Classifier

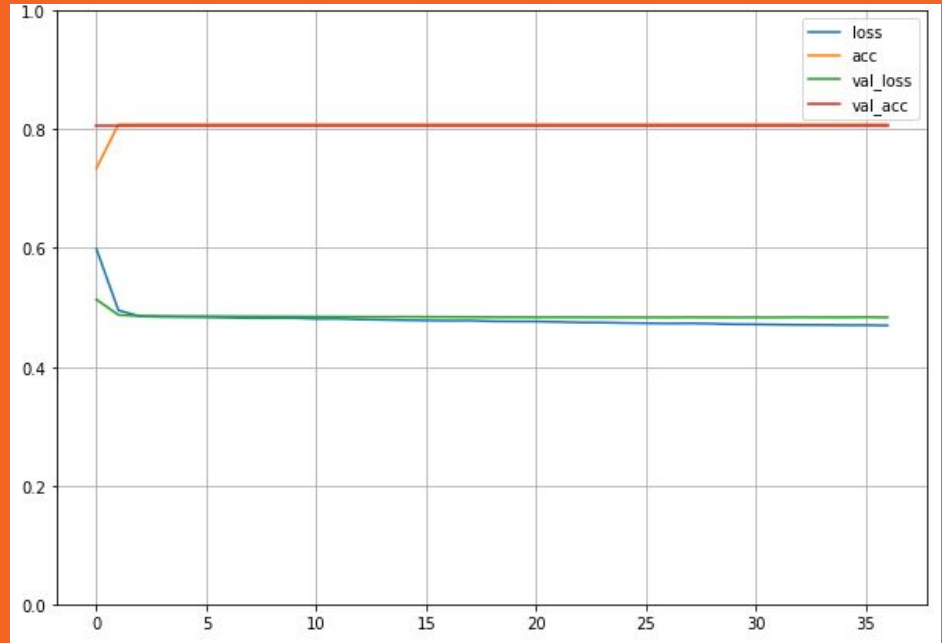
At the first trial we used the same optimizer used in extractor but there was overfitting so we changed the optimizer to adam with learning rate 10^{-7} and it was much better



Training Regressors For Contrast

Regressor

At first we used the adam optimizer with learning rate 10^{-7} but result was not good so we tried to increase the learning rate and make it 10^{-2} and it was much better



Contrast Model Statistics



Regressors

We judge the model by two methods

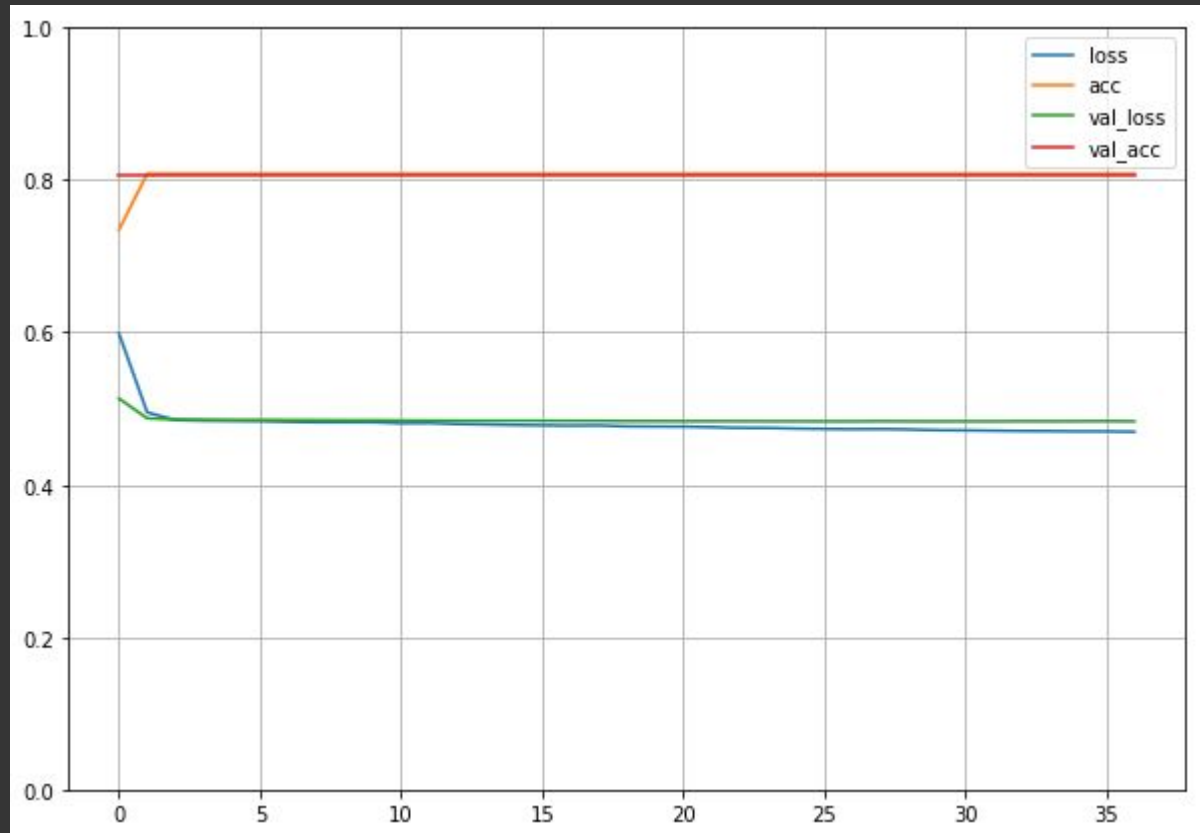
- Loss
- Accuracy

Contrast

- **Abnormal**

Loss : 0.5096

Accuracy : 0.7917

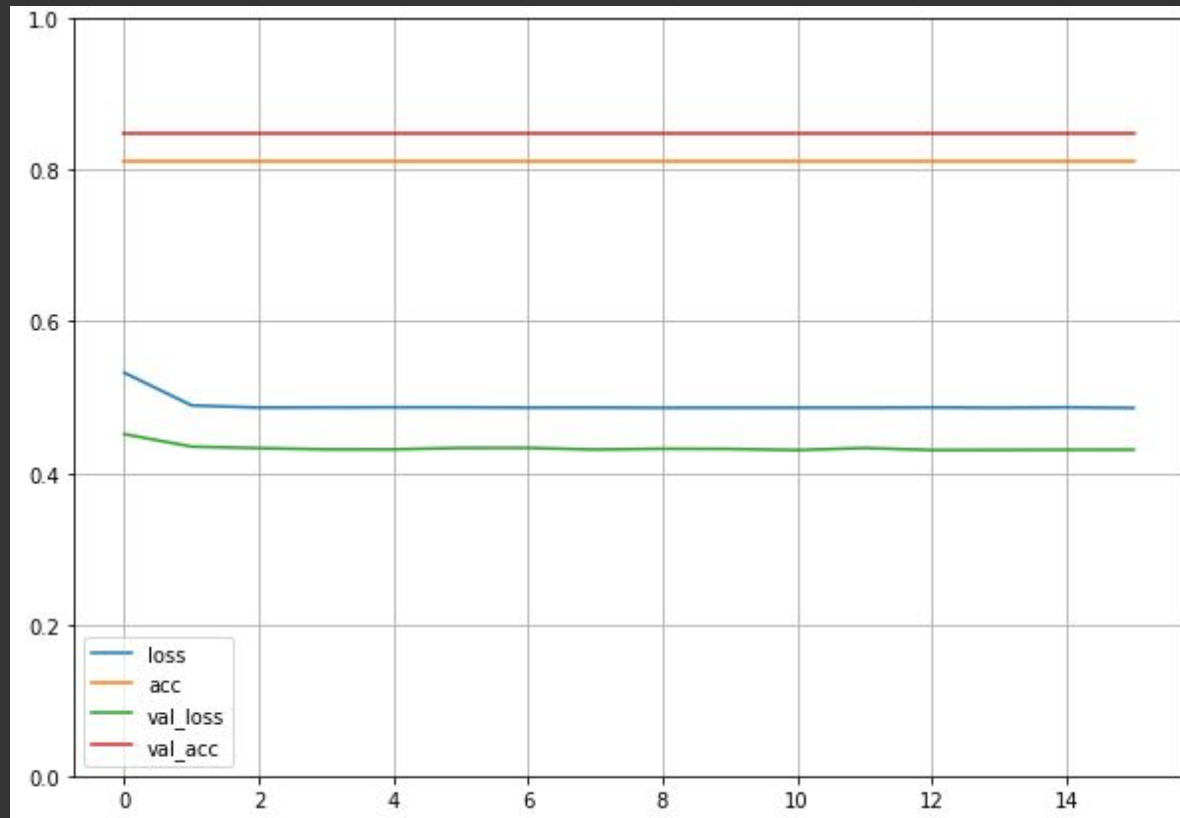


Contrast

- **ACL**

Loss : 0.8062

Accuracy : 0.5500

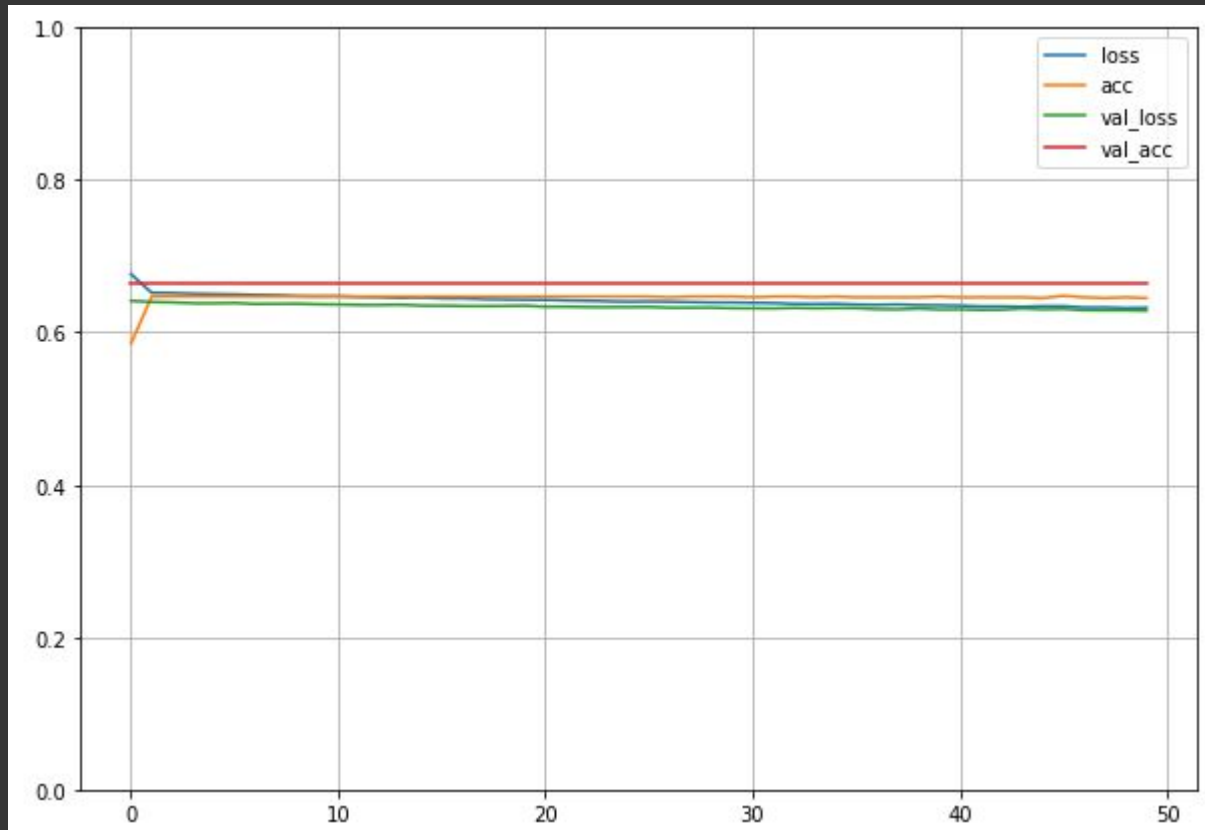


Contrast

- **Meniscal**

Loss : 0.6925

Accuracy : 0.5667



Contrast model summary

	Accuracy	Loss
Abnormal	0.7917	0.5096
ACL	0.5500	0.8062
Meniscal	0.5667	0.6925